

Program Adverbs and Tlön Embeddings

YAO LI, University of Pennsylvania, USA

STEPHANIE WEIRICH, University of Pennsylvania, USA

Free monads (and their variants) have become a popular general-purpose tool for representing the semantics of effectful programs in proof assistants. These data structures support the compositional definition of semantics parameterized by uninterpreted events, while admitting a rich equational theory of equivalence. But monads are not the only way to structure effectful computation, why should we limit ourselves?

In this paper, inspired by applicative functors, selective functors, and other structures, we define a collection of data structures and theories, which we call *program adverbs*, that capture a variety of computational patterns. Program adverbs are themselves composable, allowing them to be used to specify the semantics of languages with multiple computation patterns. We use program adverbs as the basis for a new class of semantic embeddings called *Tlön embeddings*. Compared with embeddings based on free monads, Tlön embeddings allow more flexibility in computational modeling of effects, while retaining more information about the program's syntactic structure.

ACM Reference Format:

Yao Li and Stephanie Weirich. 2022. Program Adverbs and Tlön Embeddings. *Proc. ACM Program. Lang.* 6, ICFP (July 2022), 30 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Suppose that you want to formally verify a program written in your favorite language—be it Verilog, Haskell, or C—your first step would be to translate that program and a description of its semantics to a formal reasoning system, such as Coq [Coq development team 2022]. This step is known as *semantic embedding* [Boulton et al. 1992].

There are multiple approaches to semantic embeddings. The two most well-known were proposed by Boulton et al. [1992]: *shallow embeddings*, which represent terms of the embedded language using equivalent terms of the embedding language, and *deep embeddings*, which represent terms using abstract syntax trees (ASTs) and represent their semantics via some interpretation function.

Shallow embeddings are convenient because they are simple, but they have their limitations. It is impossible to use them to state and reason about properties related to syntax, because they do not retain the syntactic structure of the original program. Furthermore, shallow embeddings fix a single semantics, so they are less robust to changes in program interpretations. Such edits require changing the translation process, in addition to the semantic domain (*i.e.*, the type used for representing the semantics of the embedded language).

On the other hand, deep embeddings are more modular thanks to an extra layer (*i.e.*, the AST) that defines the syntax of the embedded language. When we need to change the semantics, we only need to change the *interpretation* that maps the AST into some semantic domain—the translation to the formal reasoning system remains unchanged. Furthermore, the AST makes it possible to

Authors' addresses: Yao Li, hnkfliyao@gmail.com, Computer and Information Science, University of Pennsylvania, 3330 Walnut St, Philadelphia, PA, 19104, USA; Stephanie Weirich, sweirich@cis.upenn.edu, Computer and Information Science, University of Pennsylvania, 3330 Walnut St, Philadelphia, PA, 19104, USA.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

© 2022 Copyright held by the owner/author(s).

2475-1421/2022/7-ART

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

state and prove properties related to the original program’s syntactic structure. The downside is that interpreting and reasoning about properties based on this AST takes more effort than with shallow embeddings.

The pros and cons make it hard to choose between shallow and deep embeddings. Fortunately, we don’t need to commit to a single option. We can use *mixed embeddings*, a style of embedding that includes characteristics of each. In this style, parts of a language are embedded “shallowly” while other parts are embedded “deeply”. However, in any mixed embedding, we must ask: where should we draw the line to separate the shallowly embedded part from the deeply embedded part?

Recent efforts have focused on mixed embeddings based on freer monads [Kiselyov and Ishii 2015] or their variants [Dylus et al. 2019; McBride 2015; Swamy et al. 2020; Xia et al. 2020]. The style has been shown useful for representing and reasoning about effectful computation in various applications [Chlipala 2021; Christiansen et al. 2019; Foster et al. 2021; Letan et al. 2021; Nigron and Dagand 2021; Zakowski et al. 2021; Zhang et al. 2021]. Beyond these applications, this style points out a useful guideline for answering the question above. That is: modeling the pure parts of the computation “shallowly” and the effectful parts “deeply”.

Our work builds on this idea of separating pure and effectful parts in a mixed embedding, but inspects the following question: Why freer monads? We find that this is because freer monads model *one* general computation pattern that is common in many languages. However, the finding also implies that there are other computation patterns not captured by freer monads.

Following this observation, we propose a new class of mixed embeddings called *Tlön embeddings*.¹ Tlön embeddings model programs using structures called *program adverbs*, which are reifications of familiar type classes (e.g., *Applicative*, *Selective*, *Monad*) paired with equational theories. Like freer monads, these free structures can be used to combine shallowly embedded pure computation with deeply embedded computational effects. However, program adverbs provide choices in the semantics through the selection of the structure and equational theory. For example, the “statically” adverb, based on applicative functors and their free theory, models computation where control flow and data flow in the semantics are fixed. Or, by modifying the equational theory of the free applicative structure to include commutativity, we can describe computation that is “statically and in parallel”.

We make the following contributions:

- We compare the trade-offs of different styles of semantic embeddings in the context of formal reasoning and propose Tlön embeddings (Section 2).
- We define program adverbs and show how to define their syntactic parts and their semantic parts (Section 3).
- We refactor program adverbs to support composition and extension. We motivate why we want to compose program adverbs and define a composition algebra (Section 4).
- We implement composable program adverbs using the Coq proof assistant. A major challenge for implementing them in Coq is supporting extensible inductive data types [Wadler 1998]. We show one way of addressing this challenge by adapting the *Meta Theory à la Carte* (MTC) approach [Delaware et al. 2013] (Section 4).
- We identify five basic program adverbs from commonly used type classes in Haskell and we prove that these program adverbs are sound (Section 3). We also identify two add-on program adverbs that are used in combination with basic program adverbs (Section 4).

¹The name Tlön embedding is a reference to the short story *Tlön, Uqbar, Orbis Tertius* by Jorge Luis Borges. In the short story, Tlön is an imaginary world, where its parent language does not have any nouns, but only “impersonal verbs, modified by monosyllabic suffixes (or prefixes) with an adverbial value” [Borges 1940].

<i>literals</i>	b	$::= \text{true} \mid \text{false}$
<i>terms</i>	t, u	$::= x \mid b \mid \neg t \mid t \wedge u \mid t \vee u$

Fig. 1. The syntax of \mathcal{B} .

- We demonstrate the usefulness of program adverbs via three distinct language examples including a simple circuit language (Section 2), Haxl [Marlow et al. 2014], and a networked server adapted from Koh et al. [2019] (Section 5).

Additionally, we discuss the choice of adverb data types we use and alternative approaches to implement composable program adverbs in Section 6 and the related work in Section 7. We provide the Coq formalization of all the key concepts, theorems, and examples shown in this paper in our supplementary artifact [Li and Weirich 2022].

2 SEMANTIC EMBEDDINGS

In this section, we first demonstrate different forms of semantic embeddings using a simple circuit language called \mathcal{B} and compare how each form of embedding can be used to reason about programs written in this language. To distinguish the embedded language and the embedding language, we use mathematical notation to describe \mathcal{B} and use Coq code to describe its embeddings.

The syntax of \mathcal{B} appears in Fig. 1. Semantically, we want the Boolean operators to have their usual semantics. However, \mathcal{B} can read from the variables that represent references to external devices and we don't want to fix those values in the semantics. Furthermore, we don't know if the values are immutable: they might change over time, or they might change after each read, *etc.*

The four embeddings that we consider in this section are defined in the right column of Fig. 2. We use $\llbracket \cdot \rrbracket_S$, $\llbracket \cdot \rrbracket_D$, $\llbracket \cdot \rrbracket_M$, and $\llbracket \cdot \rrbracket_A$ to represent the translation from a term of \mathcal{B} to shallow, deep, and two mixed embeddings, respectively. These translations refer to the definitions in the left column as well as to the standard classes and notations for functors, monads, *etc.*, shown in Fig. 3.

To compare embeddings, we will use each to consider the following questions regarding the semantics of \mathcal{B} :

- (1) Is x equivalent to $x \wedge x$?
- (2) Is x equivalent to $x \wedge \text{true}$?
- (3) Is $t \wedge u$ equivalent to $u \wedge t$?
- (4) Is the number of variable accesses always less than or equal to 2 to the power of the circuit's depth?

Because we are modeling a circuit language that uses unknown external devices, we don't want to be able to prove or disprove property (1). This property may hold or not hold depending on the situation. If the external devices are immutable, this property will be true. Otherwise, we may be able to falsify it. In contrast, we would like our semantic embedding to give us tools to verify properties (2) and (3) because these properties should hold regardless of the properties of our external device. The former holds because on both sides of the equivalence relation we have only accessed the variable x once. The latter is a property of circuits in general: it says that the operands of \wedge are computed in parallel. The last property (4) relates a dynamic property of the semantics (the number of variable accesses) to a syntactic property of the circuit (the size of the circuit itself).

2.1 A Shallow Embedding

To use a shallow embedding to represent the semantics of \mathcal{B} , we need a way to represent the effects of reading from external devices—the most common way of doing this is using *monads* (Fig. 3). But

Shallow Embedding

Definition $\text{Reader } (A : \text{Type}) : \text{Type} :=$
 $(\text{var} \rightarrow \text{bool}) \rightarrow A.$
Definition $\text{ret } \{A\} (a : A) : \text{Reader } A :=$
 $\text{fun } _ \Rightarrow a.$
Definition $\text{bind } \{A B\} (m : \text{Reader } A)$
 $(k : A \rightarrow \text{Reader } B) : \text{Reader } B :=$
 $\text{fun } v \Rightarrow k (m v) v.$
Definition $\text{ask } (k : \text{var}) : \text{Reader } \text{bool} :=$
 $\text{fun } m \Rightarrow m k.$

$\llbracket \cdot \rrbracket_S : \text{Reader } \text{bool}$
 $\llbracket \text{true} \rrbracket_S = \text{ret true}$
 $\llbracket \text{false} \rrbracket_S = \text{ret false}$
 $\llbracket x \rrbracket_S = \text{ask } x$
 $\llbracket \neg t \rrbracket_S = \text{negb } \langle \$ \rangle \llbracket t \rrbracket_S$
 $\llbracket t \wedge u \rrbracket_S = t' \leftarrow \llbracket t \rrbracket_S; u' \leftarrow \llbracket u \rrbracket_S;$
 $\text{ret } (\text{andb } t' u')$
 $\llbracket t \vee u \rrbracket_S = t' \leftarrow \llbracket t \rrbracket_S; u' \leftarrow \llbracket u \rrbracket_S;$
 $\text{ret } (\text{orb } t' u')$

Deep Embedding

Inductive $\text{term} :=$
 $| \text{Var } (v : \text{var})$
 $| \text{Lit } (b : \text{bool})$
 $| \text{Neg } (t : \text{term})$
 $| \text{And } (t : \text{term}) (u : \text{term})$
 $| \text{Or } (t : \text{term}) (u : \text{term}).$

$\llbracket \cdot \rrbracket_D : \text{term}$
 $\llbracket \text{true} \rrbracket_D = \text{Lit true}$
 $\llbracket \text{false} \rrbracket_D = \text{Lit false}$
 $\llbracket x \rrbracket_D = \text{Var } x$
 $\llbracket \neg t \rrbracket_D = \text{Neg } \llbracket t \rrbracket_D$
 $\llbracket t \wedge u \rrbracket_D = \text{And } \llbracket t \rrbracket_D \llbracket u \rrbracket_D$
 $\llbracket t \vee u \rrbracket_D = \text{Or } \llbracket t \rrbracket_D \llbracket u \rrbracket_D$

Freer Monad Embedding

Inductive $\text{FreerMonad } (E : \text{Type} \rightarrow \text{Type}) R :=$
 $| \text{Ret } (r : R)$
 $| \text{Bind } \{X\} (m : E X)$
 $(k : X \rightarrow \text{FreerMonad } E R).$
Fixpoint $\text{bind } \{E A B\} (m : \text{FreerMonad } E A)$
 $(k : A \rightarrow \text{FreerMonad } E B) : \text{FreerMonad } E B :=$
 $\text{match } m \text{ with}$
 $| \text{Ret } r \Rightarrow k r$
 $| \text{Bind } m' k' \Rightarrow$
 $\text{Bind } m' (\text{fun } a \Rightarrow \text{bind } (k' a) k) \text{ end.}$
Variant $\text{DataEff } : \text{Type} \rightarrow \text{Type} :=$
 $| \text{GetData } (v : \text{var}) : \text{DataEff } \text{bool}.$

$\llbracket \cdot \rrbracket_M : \text{FreerMonad } \text{DataEff } \text{bool}$
 $\llbracket \text{true} \rrbracket_M = \text{Ret true}$
 $\llbracket \text{false} \rrbracket_M = \text{Ret false}$
 $\llbracket x \rrbracket_M = \text{Bind } (\text{GetData } x) \text{ Ret}$
 $\llbracket \neg t \rrbracket_M = \text{negb } \langle \$ \rangle \llbracket t \rrbracket_M$
 $\llbracket t \wedge u \rrbracket_M = t' \leftarrow \llbracket t \rrbracket_M; u' \leftarrow \llbracket u \rrbracket_M;$
 $\text{Ret } (\text{andb } t' u')$
 $\llbracket t \vee u \rrbracket_M = t' \leftarrow \llbracket t \rrbracket_M; u' \leftarrow \llbracket u \rrbracket_M;$
 $\text{Ret } (\text{orb } t' u')$

Reified Applicative Embedding

Inductive $\text{ReifiedApp } (E : \text{Type} \rightarrow \text{Type}) R :=$
 $| \text{EmbedA } (e : E R)$
 $| \text{Pure } (r : R)$
 $| \text{LiftA2 } \{X Y\} (f : X \rightarrow Y \rightarrow R)$
 $(a : \text{ReifiedApp } E X) (b : \text{ReifiedApp } E Y).$

$\llbracket \cdot \rrbracket_A : \text{ReifiedApp } \text{DataEff } \text{bool}$
 $\llbracket \text{true} \rrbracket_A = \text{Pure true}$
 $\llbracket \text{false} \rrbracket_A = \text{Pure false}$
 $\llbracket x \rrbracket_A = \text{EmbedA } (\text{GetData } x)$
 $\llbracket \neg t \rrbracket_A = \text{negb } \langle \$ \rangle \llbracket t \rrbracket_A$
 $\llbracket t \wedge u \rrbracket_A = \text{LiftA2 } \text{andb } \llbracket t \rrbracket_A \llbracket u \rrbracket_A$
 $\llbracket t \vee u \rrbracket_A = \text{LiftA2 } \text{orb } \llbracket t \rrbracket_A \llbracket u \rrbracket_A$

Fig. 2. Semantic embeddings of \mathcal{B} in Coq. We use the infix operator $\langle \$ \rangle$ to represent a functor's `fmap` method and a notation similar to Haskell's `do` notation to represent monadic binds. The functions `negb`, `andb`, and `orb` are Coq's functions defined on the `bool` type.

```

197 Class Functor (F : Type -> Type) :=
198   { fmap : forall {A B}, (A -> B) -> F A -> F B }.
199 Class Applicative (F : Type -> Type) `{Functor F} :=
200   { pure   : forall {A}, A -> F A ;
201     liftA2 : forall {A B C}, (A -> B -> C) -> F A -> F B -> F C }.
202 Class Selective (F : Type -> Type) `{Applicative F} :=
203   { selectBy : forall {A B C}, (A -> ((B -> C) + C)) -> F A -> F B -> F C }.
204 Class Monad (F : Type -> Type) `{Applicative F} :=
205   { ret : forall {A}, A -> F A ;
206     bind : forall {A B}, F A -> (A -> F B) -> F B }.
207
208 Default fmap definitions
209 Definition fmap_monad {m} `{Monad m} {a b} (f : a -> b) (x : m a) : m b :=
210   x >=> (fun y => ret (f y)).
211 Definition fmap_ap {t} `{Applicative t} {a b} (f : a -> b) (x : t a) : t b :=
212   liftA2 id (pure f) x.
213
214
215

```

Fig. 3. Coq type classes for functors, applicative functors [McBride and Paterson 2008], selective functors [Mokhov et al. 2019], and monads [Moggi 1991; Wadler 1992], as well as default definitions of fmap.

which one? A simple option is the reader monad. We show core definitions of a specialized reader monad at the top left of Fig. 2.² The translation from \mathcal{B} to `Reader bool` is given in the same figure. Following the terminology used by Svenningsson and Axelsson [2012], we call `Reader bool` the *semantic domain* of our shallow embedding. Of course, the reader monad is just one possible semantic domain, other candidates include Dijkstra monads [Swamy et al. 2013], predicate transformer semantics [Swierstra and Baanen 2019], etc.

Using the reader monad, we can prove that property (1) is true, using (\approx_S), the point-wise equality of functions. More specifically, we can prove the following Coq theorem:

```
forall x, ask x  $\approx_S$  x1 <- ask x; x2 <- ask x; ret (andb x1 x2)
```

We “ask” twice on the right hand side of the equivalence to model accessing variable x twice during program runtime. However, $x1$ equals to $x2$ in our case since nothing has changed the global store. After proving that, the theorem can be proved via a case analysis on $x1$.

However, note that our proof relies on “nothing has changed the global store,” but we don’t know if this is true, as we don’t know anything about the characteristics of the external device. Indeed, property (1) should *not* be true if we have a device where its values change over time: the value of x might change between two variable access. This is a problem with our choice of semantic domain. By choosing the reader monad, we introduce more assumptions over the semantics of \mathcal{B} , which results in proving a property that is not supposed to be true in the original language \mathcal{B} .

Although this is not a problem with the approach of shallow embedding—we can choose a different monad than the reader monad, the style does force us to choose a concrete semantic domain early. In practice, we sometimes need to change the semantic domain, either because we made a wrong assumption or because the language evolves. With shallow embeddings, we would need to change the entire translation process to change this domain.

²For simplicity, we specialize the monad so that its environment has type `var -> bool`. The commonly used reader monad is more general that the type of its environment is parameterized.

Unlike property (1), property (2) is true even though we don't know anything about the external device. This is because on both sides of the equivalence relation we have only accessed the variable x once. Property (2) can be stated as follow with our shallow embedding:

```
forall x, ask x  $\approx_S$  x1 <- ask x; ret (andb x1 true)
```

The proof follows from the theories of Coq's `bool` type and the Reader monad. However, even though this property should be true regardless of the external device, our mechanical proof still relies on the assumption that the external device is immutable—this is again because the property is stated in terms of the reader monad. If we change the shallow embedding to use a different semantic domain, we would need to prove this property again.

Property (3) is true and we can prove it to be true using our shallow embedding, but that is just a lucky hit. Even though we know nothing about the external device, there is a bisimulation between $t \wedge u$ and $u \wedge t$ because the two operands t and u run in parallel in a circuit. A proof based on our shallow embedding would, on the other hand, be based on the wrong assumption that the external device is immutable.

We cannot state property (4) with our shallow embedding. Our shallow embedding does not retain the syntactic structure of the original program so we cannot define a function that calculates the depth of the circuit.

2.2 A Deep Embedding

In a deep embedding, we first define an abstract syntax tree (AST) for \mathcal{B} . For example, we can use the term data type shown in Fig. 2. Our translation from \mathcal{B} to the term is shown in the same figure. Note that the term data type does not encode any semantic meaning.

Without an interpretation, we cannot prove any of the first three properties. This is actually ideal for answering question (1) since we know nothing about the external device so we should not be able to prove it (nor should we be able to prove it wrong!). However, by leaving the entire syntax tree uninterpreted we are now unable to prove property (2) or (3), either.

A way out of this quandary is to define a coarser *equivalence relation* for ASTs and use that relation in the statement of properties (2) and (3). For example, we can interpret each term using the reader monad (as in the shallow embedding) and use the point-wise equivalence relation for that type. The proofs are essentially the same as the above.

One advantage of the deep embedding in this case is that, if we would like to change our definition of equivalence, we can do so by choosing a different *interpretation* without changing the translation process. In other words, deep embeddings achieve better modularity by introducing an intermediate layer. The price, however, is that it takes effort to build that extra intermediate layer. This extra effort seems small here, but can become tedious with some languages, e.g., those with features like “let” that introduce variable bindings [Aydemir et al. 2005].

However, we still face a similar problem with the shallow embedding: If we would like to change the interpretation in our definition of equivalence, we need to prove our properties again. This suggests that another intermediate layer between deep and shallow embeddings might be helpful, as we will see in the next subsection.

The primary benefit we have by using the deep embedding is that we can now state and prove property (4). This is because the deep embedding gives us a representation of the program's original syntactic structure. This allows us to define the following function that counts the depth of a circuit:

```
Fixpoint depth (t : term) : nat :=
  match t with
  | Var _ => 0
  | Lit _ => 0
```



```

LEFT IDENTITY   : ret a >>= h = h a
RIGHT IDENTITY  : m >>= ret = m
ASSOCIATIVITY   : (m >>= g) >>= h = m >>= (fun x => g x >>= h)

```

Fig. 4. The monad laws. The $\gg=$ symbol is the infix operator for bind. Proving these laws for `FreerMonad` in Coq relies on the axiom of functional extensionality.

```

| Neg t => depth t + 1
| And t u => max (depth t) (depth u) + 1
| Or t u => max (depth t) (depth u) + 1
end.

```

Since we assume a straightforward semantics for \mathcal{B} , the number of variable access at runtime equals to the number of variables appeared in a term, so we can directly prove property (4) by an induction over the term data type.

2.3 A Mixed Embedding Based on Freer Monads

A semantic embedding can be partially shallow and partially deep. We use the term *mixed embeddings* to describe embeddings with this property. One style of mixed embeddings that is popular today is based on *freer monads* [Chlipala 2021; Dylus et al. 2019; Letan et al. 2021; McBride 2015; Nigron and Dagand 2021; Swamy et al. 2020; Xia et al. 2020]. In this type of mixed embeddings, the pure parts of the program are embedded shallowly, while effects are embedded deeply (and abstractly) using algebraic data types “connected” by freer monads.

The core definitions of freer monads are in the left column of Fig. 2. The `FreerMonad` data type is parameterized by an abstract effect E of `Type -> Type` and a return type R of `Type`. Conceptually, it collects all the deeply embedded effects E in a right-associative monadic structure.

For any effect type, `FreerMonad E` is a monad as demonstrated by the `Ret` constructor and `bind` function. The `bind` function pattern matches its first argument m and, in the case of `Bind`, passes its second arguments k to the continuation of m . This “smart constructor” ensures that binds always associate to the right.

To embed \mathcal{B} , we model reading data from external devices using the effect type `DataEff`. This datatype includes only one (abstract) effect, called `GetData`. This constructor represents a data retrieval with the variable $v : \text{var}$ that returns an unknown `bool`. Similar to how the term data type says nothing about the semantics of \mathcal{B} , the effect data type `DataEff` says nothing about the semantics of a data read. As a result, we say that the effects are embedded deeply in this style.

The embedding function appears on the right side of Fig. 2. The translation strategy is almost the same as embedding \mathcal{B} using the reader monad. The only exception is in the variable case (the effectful part): here the `Bind` constructor marks the occurrence of the `GetData` effect.

In this mixed embedding, the pure parts of a \mathcal{B} program have been translated to a shallow semantic domain, but the effectful parts remain abstract. It turns out that this separation is useful for both questions (1) and (2).

For question (1), we cannot answer it. This is desirable since we don’t know if it’s true without knowing more about the external device.

We can prove that property (2) is true even though the read effect is not interpreted—this is because the property follows from the monad laws (Fig. 4). However, we cannot prove property (3) because the commutativity law is not one of the monad laws.

Ideally, we would also like to state and prove property (4). However, the dynamic nature of freer monads forbids us from statically inspecting the syntactic structure of the program. Interpreting the embedding does not help us, either, since that would not preserve the original syntactic structure.

Our success with questions (1) and (2) suggests that we have found an useful intermediate layer between shallow and deep embeddings, but our failure in stating or proving properties (3) and (4) indicates that we haven't yet found the right representation.

2.4 Another Mixed Embedding Based on Reified Applicative Functors

The last embedding shown in the figure uses a type that reifies the interface of *applicative functors* (Fig. 3). As in freer monads, this datatype is parameterized by deeply embedded abstract effects. These effects, of type $E \rightarrow R$, are recorded by the `EmbedA` data constructor.

However, instead of constructors for `ret` and `bind`, this datatype includes constructors for `pure` and `liftA2`, the two operations that define applicative functors.³ The `Pure` constructor shallowly “embeds” a pure computation into the domain, and `LiftA2` “connects” two computations that potentially contain effect invocations. These constructors provide a trivial implementation of the `Applicative` type class for this datatype.

The translation of \mathcal{B} to this datatype uses a deep embedding of variable reads, using the `EmbedA` data constructor with the `DataEff` type from the previous embedding. Because, as in freer monads, this effect is modeled abstractly, we cannot prove or disprove (1).

The translation function uses the applicative interface in the datatype to translate the constants, unary and binary operators. These components are modeled shallowly (*i.e.*, as Boolean constants and operators), but the program's syntactic structure is retained by the translation. However, because of the retainment, we need an additional equivalence relation to equate semantically equivalent terms that are not syntactically equal. To prove (2), we include the right identity law of applicative functors in the equivalence (denoted by \cong):

$$\frac{\forall y, (\text{fun } _ \ x \Rightarrow x) \ a \ y = f \ a \ y}{\text{liftA2 } f \ (\text{pure } a) \ b \cong b}$$

This law is sufficient to show that (2) holds.

To model the parallelism of circuits, we can include the commutativity law in the equivalence:

$$\text{liftA2 } f \ a \ b \cong \text{liftA2 } (f \text{ flip }) \ b \ a$$

This is sufficient to show (3). Note that this is not one of the applicative functor laws. We defer showing the soundness of including this rule in the equivalence to Section 3.4.

This embedding also preserves enough of the syntax of the original program to prove (4). To do so, we must first calculate the depth of circuits and the number of variables under this encoding.

```

Fixpoint app_depth {E A} (t : ReifiedApp E A) : nat :=
  match t with
  | EmbedA _ => 0
  | Pure _ => 0
  | LiftA2 _ t u => 1 + max (app_depth t) (app_depth u)
  end.

```

We omit the function that counts the number of variables as it is similar to `app_depth`. Then we can formalize (4) in Coq as follow:

³Alternatively, `Applicative` can also be defined by `pure` and another operation `<*>` of type $F \ (A \rightarrow B) \rightarrow F \ A \rightarrow F \ B$, where F is an `Applicative` instance. These two definition are equivalent, as we can derive the definition of `<*>` from `liftA2` and vice versa.

Theorem `heightAndVar : forall (c : ReifiedApp DataEff bool),
app_numVar c <= Nat.pow 2 (app_depth c).`

The theorem is provable by an induction over c .

Furthermore, this embedding also allows us to reason about semantical properties that depend on syntactic structures of circuits. One example is a semantics with some cost model. In the semantics, we may not want our equivalence to equate, for example, $x \wedge y \wedge z \wedge w$ and $(x \wedge y) \wedge (z \wedge w)$ because they are not equivalent in their costs when parallelization is present. Indeed, we cannot show that they are equivalent with our embedding due to the absence of associativity in our equivalence.

2.5 Tlön embeddings

Just as the reader monad models *one* particular effect, freer monads model *one* particular computation pattern. Unfortunately, that particular computation pattern is not suitable for our \mathcal{B} example, because it does not model parallel computation (*i.e.*, property (3)), nor does it capture the static data and control flows (*i.e.*, property (4)). Instead we saw that the mixed embedding in the previous subsection, based on reified applicative functors, is a better approach.

Can we generalize the key idea even further? If we go beyond \mathcal{B} , we might need to model other computation patterns. Are there other mixed embeddings that would be suitable for these tasks? How might we derive them?

To that end, we identify a novel set of mixed embeddings that we call *Tlön embeddings*. The goal of these embeddings is to provide flexibility in our models of effectful computation.⁴ We define Tlön embeddings by identifying a set of *program adverbs* that specify the embedding type and equational theory used in the embedding. For example, the embedding in Section 2.4 is based on an adverb composed of the `ReifiedApp` type and *some* rules of commutative applicative functors.

The flexibility that program adverbs provide can perhaps be understood by comparing them with effects: effects *do* certain actions, and program adverbs model *how* these actions are done—similar to the difference between verbs and adverbs. For example, the adverb we used in Section 2.4 is called “statically and in parallel”, which states that there is a static dependency between different effect invocations and some of these effect invocations are executed in parallel.

In the next section, we define our set of program adverbs more precisely and discuss the reasoning principles that they provide for effectful computation.

3 PROGRAM ADVERBS

Program adverbs are the building blocks of Tlön embeddings. Mathematically, they are composed of two parts: a syntactic part, called the adverb data type, and a semantic part, called the adverb theory. More formally, we define program adverbs as follow:⁵

Definition 3.1 (Program Adverb). A program adverb is a pair (D, \cong_D) . D is called the adverb data type and is parameterized by an effect E and a return type R . The \cong_D operation is called the adverb theory of D . It is a binary operation that defines a bisimulation relation on $D(E, R)$ for any E and R .

In the rest of the paper, we abbreviate \cong_D as \cong when D is clear from the context.

⁴Here, we define effects as communications with external environment that are performed by some explicit operations. For example, *mutable states* are effects which can be explicitly incurred by operations such as `get` and `set`. For the same reason, we also consider I/O (with operations like `read`, `print`, *etc.*), exceptions (with operations like `throw`, *etc.*) as effects.

⁵The Coq code of all definitions and theorems shown in this section can also be found in our supplementary artifact.

```

442  (* Streamingly *)
443  Inductive ReifiedFunctor (E : Type -> Type) (R : Type) : Type :=
444  | EmbedF (e : E R)
445  | FMap {X : Type} (g : X -> R) (f : ReifiedFunctor E X).
446
447  (* Statically and StaticallyInParallel *)
448  Inductive ReifiedApp (E : Type -> Type) (R : Type) : Type :=
449  | EmbedA (e : E R)
450  | Pure (r : R)
451  | LiftA2 {X Y : Type} (f : X -> Y -> R)
452    (a : ReifiedApp E X) (b : ReifiedApp E Y).
453
454  (* Conditionally *)
455  Inductive ReifiedSelective (E : Type -> Type) (R : Type) : Type :=
456  | EmbedS (e : E R)
457  | PureS (r : R)
458  | SelectBy {X Y : Type} (f : X -> ((Y -> R) + R))
459    (a : ReifiedSelective E X) (b : ReifiedSelective E Y).
460
461  (* Dynamically *)
462  Inductive ReifiedMonad (E : Type -> Type) (R : Type) : Type :=
463  | EmbedM (e : E R)
464  | Ret (r : R)
465  | Bind {X : Type} (m : ReifiedMonad E X) (k : X -> ReifiedMonad E R).

```

Fig. 5. The adverb data types

In Coq terms, an adverb data type D has the type $(\text{Type} \rightarrow \text{Type}) \rightarrow \text{Type} \rightarrow \text{Type}$. The first parameter of $\text{Type} \rightarrow \text{Type}$ is the effect E and it's parameterized by its own return type; the second parameter is the return type R . The adverb theory \cong is a typed binary relation.⁶ More concretely:

```

472  Class Adverb (D : (Type -> Type) -> Type -> Type) :=
473  { Bisim {E R} : relation (D E R) ;
474    equiv {E R} : Equivalence (@Bisim E R) }.
475  Notation "a ≅ b" := (Bisim a b).

```

where D is the adverb data type, Bisim is the adverb theory \cong , and equiv is a proof showing that Bisim is an equivalence relation. The datatype relation is defined as:

```

479  Definition relation (A : Type) := A -> A -> Prop.

```

This definition is overly general, so we focus our attention only on program adverbs that are *sound* according to the definition that we will develop below. Furthermore, in this paper we will only consider adverbs defined by reifying classes of functors.

3.1 Adverb Data Types and Theories

The four key adverb data types, shown in Fig. 5, are derived from the four type classes shown in Fig. 3. We have already seen one before in the applicative embedding in Fig. 2. Other definitions

⁶In addition to bisimulation relations, we can also define refinement relations on program adverbs. We will show in Section 4.3 some adverbs with refinement relations, but bisimulation relations would suffice for most adverbs, so we only include them in the core definitions of adverb theories. Refinement relations can be added on demand.

Congruence Rule

$$\text{CONGRUENCE} : \frac{a1 \cong a2 \quad b1 \cong b2}{\text{liftA2 } f \ a1 \ b1 \cong \text{liftA2 } f \ a2 \ b2}$$

Applicative Functor Laws

$$\text{LEFT IDENTITY} : \frac{\forall y, (\text{fun } _ \ x \Rightarrow x) \ a \ y = f \ a \ y}{\text{liftA2 } f \ (\text{pure } a) \ b \cong b}$$

$$\text{RIGHT IDENTITY} : \frac{\forall x, (\text{fun } x \ _ \Rightarrow x) \ x \ b = f \ x \ b}{\text{liftA2 } f \ a \ (\text{pure } b) \cong a}$$

$$\text{ASSOCIATIVITY} : \frac{\forall x \ y \ z, f \ x \ y \ z = g \ y \ z \ x}{\text{liftA2 } \text{id} \ (\text{liftA2 } f \ a \ b) \ c \cong \text{liftA2 } (\text{flip } \text{id}) \ a \ (\text{liftA2 } g \ b \ c)}$$

$$\text{NATURALITY} : \frac{\forall x \ y \ z, p \ (q \ x \ y) \ z = f \ x \ (g \ y \ z)}{\text{liftA2 } p \ (\text{liftA2 } q \ a \ b) \cong \text{liftA2 } f \ a \ . \ \text{liftA2 } g \ b}$$

Equivalence Properties

$$\text{REFLEXIVITY} : \frac{}{a \cong a} \quad \text{SYMMETRY} : \frac{a \cong b}{b \cong a}$$

$$\text{TRANSITIVITY} : \frac{a \cong b \quad b \cong c}{a \cong c}$$

Fig. 6. The equivalence relations for ReifiedApp. The infix operator \cdot denotes function compositions.

follow a similar pattern: the constructors of each data type include one for embedding effects (of type $E \rightarrow R$) and a constructor that reifies the interface of each method of the type class.

In addition to an adverb data type, every program adverb also comes with some theories, defined by a bisimulation relation \cong . The purpose of the \cong relation is to equate all computations that are semantically equivalent regardless of what effects are present.

For example, an adverb called *Statically* is composed of the *ReifiedApp* datatype with an equational theory based on three sorts of rules: (1) a congruence rule with respect to *LiftA2*, (2) the laws of applicative functors [McBride and Paterson 2008], and (3) the equivalence properties (*i.e.*, reflexivity, symmetry, transitivity). We show the concrete rules in Fig. 6.

Why do we call this adverb *Statically*? The data dependency in the *LiftA2* constructor of *ReifiedApp* shows that the data type imposes a “static” data flow and control flow on the computation: we will always need to run both parameters of type *ReifiedApp* $E \rightarrow A$ and *ReifiedApp* $E \rightarrow B$ to get the result of type *ReifiedApp* $E \rightarrow C$, *i.e.*, we cannot skip either computation. In addition, neither of the two parameters depends on the result of the other, which allows us to statically inspect either of them without running the other.

Remark. The adverb data types and their associated theories form free structures similar to those in Capriotti and Kaposi [2014]; Kiselyov and Ishii [2015]; Mokhov [2019]; Mokhov et al. [2019]. However, one distinction is that we intentionally do not normalize the adverb data types to preserve syntactic structures. To distinguish un-normalized free structures and normalized free structures, we use the term *reified* structures to describe the former and the term *free* structures to exclusively

```

540 Fixpoint interpA {E I : Type -> Type} {Applicative I} {A : Type}
541   (interpE : forall A, E A -> I A) (t : ReifiedApp E A) : I A :=
542   match t with
543   | EmbedA e => interpE _ e
544   | Pure a => pure a
545   | LiftA2 f a b => liftA2 f (interpA interpE a) (interpA interpE b)
546   end.
547

```

Fig. 7. The interpretation from ReifiedApp to any instance of the Applicative type class.

describe the latter. We defer the detailed comparison and trade-offs between reified structures and free structures to [Section 6](#).

3.2 Adverb Simulation

One important property of ReifiedApp is that it can be interpreted to any other instance of the Applicative class, as long as its embedded effects can be interpreted to that instance. We can show this via the abstract interpreter `interpA` shown in [Fig. 7](#). The interpreter shows that given *any* effect E and *any* instance I of Applicative, as long as we can find an effect interpretation from E to I for any type A , we can interpret a `ReifiedApp E A` to an `I A` for any type A .

For example, we can interpret a `ReifiedApp DataEff` to the reader applicative functor ([Fig. 2](#))⁷ by supplying the following function to the parameter `interpE` of `interpA`:

```

563 Definition interpDataEff {A : Type} (e : DataEff A) : Reader A :=
564   match e with GetData v => ask v end.
565

```

Similarly, we can interpret `ReifiedApp DataEff` to other semantic domains that are applicative functors.

Why do we care if ReifiedApp can be interpreted into any instance of Applicative? This is because different instances of Applicative model different effects—if we have a data structure that can be interpreted to all instances, we can develop a theory of it that can be used for reasoning about properties that are true regardless of what effects are present.

To make the relation between an adverb data type like ReifiedApp and a class of functor like Applicative more precise, we define the following *adverb simulation* relation:

Definition 3.2 (Adverb Simulation). Given an adverb data type D , a class of functor C , and a transformer T on all instances of C , we say that there is an adverb simulation from D to C under T , written $D \models_T C$, if we can define a function that, for any effect type E , instance F of type class C , and interpreter f from $E(A)$ to $F(A)$ for any type A , interprets a value of $D(E, A)$ to $T(F)(A)$ for any type A .

We add some flexibility to this definition by making it parameterize over a transformer T —we do not need this extra flexibility for now, but we will see why it is useful in [Section 3.4](#).

We also define an *adverb interpretation* as follow:

Definition 3.3 (Adverb Interpretation). Given an adverb data type D , a class of functor C , and a transformer T on all instances of C , an interpreter I that shows $D \models_T C$ is called an adverb interpretation, and we write $I \in D \models_T C$.

⁷Every monad is also an applicative functor, so the reader monad is also a reader applicative functor.

Our `interpA` in Fig. 7 is an adverb interpretation. More specifically, we say that

$$\text{interpA} \in \text{ReifiedApp} \models_{\text{IdT}} \text{Applicative}.$$

where the `IdT` transformer is an identity `Applicative` transformer that “does nothing”. In the rest of the paper, when we have $D \models_{\text{IdT}} C$ for any D and C , we abbreviate it as $D \models C$.

3.3 Sound Adverb Theories

To know that our adverb theory is *sound*, i.e., it doesn’t equate computations that are not semantically equivalent, we define the following soundness property of adverb theories:

Definition 3.4 (Soundness of Adverb Theories). Given a program adverb (D, \cong) and an adverb interpretation $I \in D \models_T C$, we say that the adverb theory \cong is sound with respect to I if there exist a lawful equivalence relation \equiv such that for all $d_1, d_2 \in D$,

$$d_1 \cong d_2 \implies I(d_1) \equiv I(d_2).$$

Let us use `idT` for the transformer T for the moment. The equivalence relation \equiv on C is lawful if they respect the congruence laws and the class laws of C . For `Applicative`, we use the common applicative functor laws regarding \equiv . Based on the soundness of adverb theories, we can define the following soundness property of program adverbs with respect to their adverb interpretations:

Definition 3.5 (Soundness of Program Adverbs). Given a program adverb (D, \cong) and an adverb interpretation $I \in D \models_T C$, we say that the adverb is sound if the \cong relation is sound with respect to I .

We can now prove that the `Statically` adverb is sound:

THEOREM 3.6. *The `Statically` adverb $(\text{ReifiedApp}, \cong)$ is sound with respect to the adverb interpretation $\text{interpA} \in \text{ReifiedApp} \models \text{Applicative}$.*

PROOF. By induction over the \cong relation. □

3.4 “Statically and in Parallel”

Two adverbs can use the same data type yet differ in their theories. Let’s look at a variant of the `Statically` adverb called `StaticallyInParallel`. As its name suggests, it adds parallelization to a static computation pattern.

Recall that the two computations connected by `liftA2` do not depend on each other. This suggests that an implementation of `liftA2` can choose to run them in parallel. Indeed, that observation is one of the key ideas behind `Haxl` [Marlow et al. 2014].

Based on this idea, we also define the `StaticallyInParallel` adverb. The adverb data type of this adverb is the same as that of `Statically`. However, its theory differs from `Statically` in the following ways: (1) it adds the commutativity rule:

$$\text{liftA2 } f \ a \ b \cong \text{liftA2 } (\text{flip } f) \ b \ a$$

and (2) it does not include the associativity and naturality rules (Fig. 6).

The addition of commutativity rule states that the order that effects are invoked does not matter. Note that compared with other rules, the commutativity rule is not satisfied by every applicative functor. This might suggest that we should not add it to the theory, as it might be a theory that only holds for certain effects. Nevertheless, we can prove the soundness of the adverb theory with respect to the following adverb simulation:

$$\text{ReifiedApp} \models_{\text{PowerSet}} \text{Applicative}$$

```

638 Definition PowerSet (I : Type -> Type) (A : Type) := I A -> Prop.
639
640 Definition embedPowerSet {A : Type} (a : I A) : PowerSet I A := fun r => r ≡ a.
641
642 Definition purePowerSet {A : Type} (a : A) : PowerSet I A := fun r => r ≡ pure a.
643
644 Definition liftA2PowerSet {A B C} (f : A -> B -> C)
645       (a : PowerSet I A) (b : PowerSet I B) : PowerSet I C :=
646   fun r => exists a', a a' /\ exists b', b b' /\
647       (liftA2 f a' b' ≡ r \/ liftA2 (flip f) b' a' ≡ r).
648
649 Definition EqPowerSet {A} : relation (PowerSet I A) :=
650   fun p q => forall a, p a <-> q a.
651
652

```

Fig. 8. The core definitions of a powerset applicative functor transformer.

The `PowerSet` transformer is *a transformer on applicative functors* and its core definitions are shown in Fig. 8. The key of `PowerSet` is the `liftA2PowerSet` operation. When executed, it creates two nondeterministic branches (indicated by the disjunction `\|`): on one branch, it computes `a' : I A` before `b' : I B`, and vice versa on the other branch. Intuitively, this is to model the nondeterministic execution order in a parallel evaluation. Many of these operations depend on `≡`, which is the lawful `≡` relation on `I`.

LEMMA 3.7. *If \equiv is a lawful equivalence relation on `Applicative`, `EqPowerSet` is an equivalence relation on `PowerSet I` that satisfies congruence, left identity, right identity, and commutativity laws.*

PROOF. By definition. □

Note that `EqPowerSet I` does not satisfy the associativity or naturality laws. Consider that we have `liftA2PowerSet id (liftA2PowerSet f a b) c`, for some `f`, `a`, `b`, and `c`: one of the possible evaluations in this powerset is `liftA2 id (liftA2 (flip f) b a) c`, which does not belong to the powerset of `liftA2PowerSet (flip id) a (liftA2PowerSet g b c)`, for some `g` that is equivalent to `flip f`. The case for naturality is similar. For this reason, we do not include these two rules in `≡`. We do not know if there exists an alternative nontrivial transformer with an equivalence relation that satisfies *all* the applicative laws in addition to commutativity.

Nevertheless, we can show the following theorem with the help of Lemma 3.7:

THEOREM 3.8. *The adverb is sound: $\text{ReifiedApp} \models_{\text{PowerSet}} \text{Applicative}$.*

PROOF. We can construct an `interpPowerSet` $\in \text{ReifiedApp} \models_{\text{PowerSet}} \text{Applicative}$ by modifying `interpA` (Fig. 7) so that it uses `embedPowerSet` on the `EmbedA` case, `purePowerSet` on the `Pure` case, and `liftA2PowerSet` on the `LiftA2` case. With the help of Lemma 3.7, we can show that for all $d_1, d_2 \in \text{ReifiedApp}$,

$$d_1 \cong d_2 \implies \text{interpPowerSet}(d_1) \equiv \text{interpPowerSet}(d_2)$$

where `≡` is `EqPowerSet`. □

Intuitively, we can define `StaticallyInParallel` as an adverb because, even though with an effect running computations in different order might return different results, a language can be implemented in a parallel way such that the difference in evaluation orders is no longer observable.

The lack of associativity and naturality rules in the theory of `StaticallyInParallel` might initially sound limiting, but, as we have shown in the end of [Section 2.4](#), it turns out to be desirable in cases like circuits.

3.5 Other Basic Adverbs

Besides `Statically` and `StaticallyInParallel`, we also identify three other basic adverbs, namely `Streamingly`, `Conditionally`, and `Dynamically`, defined using the adverb data types in [Fig. 5](#).

Streamingly. This program adverb simulates `Functor` under `IdT`. The most simple form of stream processing computes the data directly as it is received. This is captured by the `fmap` interface ([Fig. 3](#)).

Dynamically. This adverb simulates `Monad` ([Fig. 3](#)). A monad is the most expressive and dynamic among all four classes of functors thanks to its core operation `bind`. Any kind of computation can happen in the second operand and we can't know it without knowing a value of type `A`, which we can only get by running the first operand. This program adverb is commonly used in representing many programming language for its expressiveness, but it also allows for the least amount of static reasoning.

Unlike `Statically`, this variant does not have an `InParallel` variant. This might be surprising because there are many commutative monads. However, those monads are commutative because their specific effects are commutative. We cannot define a general powerset *monad transformer* that can make any monad satisfy the commutativity law.

Conditionally. We use this adverb to model conditional execution. The definition of its adverb data type is shown in [Fig. 5](#). It reifies the `Selective` type class ([Fig. 3](#)). The signature operation of `Selective` is the `selectBy` operation. Loosely, “applying” a function of type `A -> ((B -> C) + C)` to a computation of type `F A` gets you either `F (B -> C)` or `F C`. In the first case, you will need to run the computation of type `F B`. You don't *need* to run the computation of type `F B` in the second case, but you can still choose to run it.

Because we can encode conditional execution with this adverb, it is more expressive than `Statically`. However, the extra expressiveness also makes static analysis less accurate. Since we cannot know statically if the computation `F B` in `selectBy` is executed, we can only get an under-approximation (assuming that `F B` is not executed) and an over-approximation (assuming that `F B` is executed) of the effects that would happen, but not an exact set.

Even though we derive this adverb by reifying `Selective`, we do not wish to model the adverb's theory using the laws of selective functors. This is because the laws of selective functors do not distinguish them from applicative functors. Indeed, every applicative functor is also a selective functor (by running the second argument even when not required) and vice versa, so adhering to the “default” laws do not allow us to prove more properties. Therefore, we add one simple rule to the selective functor laws:

$$\text{select } (\text{inr } \langle \$ \rangle a) b \cong a$$

The function `select` has the type `F (A + B) -> F (A -> B) -> F B`, where `F` is an instance of `Selective`. It is equivalent to

```
selectBy (fun x => match x with
  | inl x => inl (fun y => y x)
  | inr x => inr x
end).
```

This rule forces `select` to ignore the second argument when it does not need to be run. However, we can no longer show that `Conditionally` adverb simulates `Selective` by adding this laws, because

\cong is no longer an under-approximation of \equiv . Instead, we show the following adverb simulation:

$$\text{ReifiedSelective} \models \text{Monad}$$

In this way, Conditionally serves as a compromise between Statically and Dynamically. Its adverb data type is more similar to Statically and allows for some static analysis, while its theories are more similar to Dynamically.

4 COMPOSABLE PROGRAM ADVERBS

From a monad instance, we can derive an applicative functor instance. From an applicative functor instance, we can derive a functor instance. We can derive a selective instance from an applicative functor and vice versa.⁸ This subsumption hierarchy among classes of functors means that we can choose the most expressive abstract interface of a data type, and that choice automatically includes the less expressive interfaces.

However, although we can derive a “default” applicative functor from a monad, we don’t always want to do that—e.g., we may want to define different behavior for `liftA2` than the one derived from `bind`. Indeed, Haxl is one such example, where `bind` is defined as a sequential operation and `liftA2` is parallel so that certain tasks with no data dependencies can be automatically parallelized [Marlow et al. 2014]. In the program adverbs terminology, the semantics of their language is composed of a “statically and in parallel” adverb and a “dynamically” adverb.

In addition, some languages may have a subset that corresponds to the “statically” adverb and some extensions that correspond to “dynamically”. If we only use the “dynamically” adverb to reason about programs written in this language, we lose the ability to state properties for the “statically” subset.

We need a way to compose multiple program adverbs. Therefore, in this section, we refactor program adverbs to *composable program adverbs*.

4.1 Uniform Treatment of Effects and Program Adverbs

Effects are commonly considered secondary to monads. This treatment of effects carries over to the approaches based on freer monads and our previous implementation of program adverbs, where the effects are a parameter of adverb data types.

This approach works well when we use one fixed program adverb, but needs an update when multiple adverbs are involved. This is because, in both scenarios we mentioned earlier, our intention is not to combine program adverbs that each contain their own set of effects—we would like the composed program adverbs to share the same set of effects. One solution is requiring that we can only join program adverbs when they share the same set of effects, but that would require extra machinery.

In our work, we choose to give a uniform treatment to effects and program adverbs. Figure 9 shows our algebra for effects and program adverbs. The algebra includes an \oplus operator which is a *disjoint union* of effects and adverb data types. We define an equivalence relation \approx on effects and adverb data types as follows: for all A, B that are effects and adverb data types, $A \approx B$ if there exists a *bijection* between A and B . Similarly, we define an \uplus operator for the disjoint union of adverb theories. We define an equivalence relation \Leftrightarrow on adverb theories as follows: for any adverb data type D and adverb theories P, Q , which are adverb theories of D , $P \Leftrightarrow Q$ if $a P b \iff a Q b$ for all $a, b \in D$, where \iff is the logical symbol for “if and only if”. Properties of this algebra are also shown in Fig. 9.

⁸This is one special thing about selective functors: every selective functor is an applicative functor and the reverse is also true. However, separating these two classes is still useful because the automatically derived instances might not be what we want, as discussed in Mokhov et al. [2019].

<i>effects and adverb data types</i>	A, B, C	$::= \text{Effect } E \mid \text{AdverbDataType } D \mid A \oplus B$
<i>adverb theories</i>	P, Q, R	$::= \text{AdverbTheory } \approx_D \mid P \uplus Q$
Properties of \oplus		
COMMUTATIVITY :		$A \oplus B \approx B \oplus A$
ASSOCIATIVITY :		$(A \oplus B) \oplus C \approx A \oplus (B \oplus C)$
Properties of \uplus		
COMMUTATIVITY :		$P \uplus Q \Leftrightarrow Q \uplus P$
ASSOCIATIVITY :		$(P \uplus Q) \uplus R \Leftrightarrow P \uplus (Q \uplus R)$
IDEMPOTENCE :		$P \uplus P \Leftrightarrow P$

Fig. 9. The algebra for effects and composable program adverbs.

4.2 The Coq Implementation

All the adverb data types we have seen (Fig. 5) are recursive. When we compose these program adverbs, we cannot simply put these inductive types into a sum type—we need to adapt each adverb so that it recurses on the new composed adverb rather than itself. In other words, we need *extensible inductive types*. However, extensible inductive types are not directly supported by most formal reasoning systems including Coq. In fact, how to support extensible inductive types is part of an open problem known as *the expression problem* [Wadler 1998].

In this paper, we address the problem and implement composable adverbs in Coq using a technique presented in *Meta Theory à la Carte* (MTC) [Delaware et al. 2013]. The key idea of MTC is using Church encodings of data types [d. S. Oliveira 2009; Wadler 1990] instead of Coq’s native inductive types. We apply and extend this idea to define the two least fixpoint operators Fix1 and FixRel that work on adverb data types and adverb theories, respectively. We show the definitions of these operators in Fig. 10a.

We define the disjoint union \oplus by first refactoring the types of adverb data types and effects. We make both adverb data types and effects have the type $(\text{Set} \rightarrow \text{Set}) \rightarrow \text{Set} \rightarrow \text{Set}$ where the first parameter is a *recursive parameter* and the second parameter is a return type. We can then define \oplus simply as a sum type on $(\text{Set} \rightarrow \text{Set}) \rightarrow \text{Set} \rightarrow \text{Set}$, as shown in Fig. 10b. Similarly, we define \uplus as a sum type on $(\text{forall } (A : \text{Set}), \text{ relation } (F A)) \rightarrow \text{forall } (A : \text{Set}), \text{ relation } (F A)$ for any $F : \text{Set} \rightarrow \text{Set}$.

Figure 11a shows the definitions of composable adverb data types. Compared with the adverb data types in Fig. 5, a composable adverb data type replaces the effect parameter (which is named E) with a recursive parameter (which is named K) so that it “recurses” on K instead of itself.

We also factor out the Pure constructor, a common part shared by multiple basic adverb data types, as a separate composable adverb data type called ReifiedPure . In this way, we avoid introducing multiple Pure constructors, e.g., by combining Statically and Conditionally . Furthermore, we remove the Embed constructors in composable adverb data types. Thanks to the uniform treatment of effects and program adverbs, we can now embed effects simply by including them in K , so we have no need for those constructors.

As an example, we can define an “inductive type” $T : \text{Set} \rightarrow \text{Set}$ that is composed of ReifiedPure , ReifiedApp , and some effect $E : (\text{Set} \rightarrow \text{Set}) \rightarrow \text{Set} \rightarrow \text{Set}$ as follow:

Definition $T := \text{Fix1 } (\text{ReifiedPure} \oplus \text{ReifiedApp} \oplus E)$.

```

834 Definition Alg1 (F : (Set -> Set) -> Set -> Set) (E : Set -> Set) : Type :=
835   forall {A : Set}, F E A -> E A.
836 Definition Fix1 (F : (Set -> Set) -> Set -> Set) (A : Set) :=
837   forall (E : Set -> Set), Alg1 F E -> E A.
838
839 Definition AlgRel {F : Set -> Set}
840   (R : (forall (A : Set), relation (F A)) -> forall (A : Set), relation (F A))
841   (K : forall (A : Set), relation (F A)) : forall (A : Set), relation (F A) :=
842   fun A (a b : F A) => R K _ a b -> K _ a b.
843 Definition FixRel {F : Set -> Set}
844   (R : (forall (A : Set), relation (F A)) -> forall (A : Set), relation (F A))
845   : forall (A : Set), relation (F A) :=
846   fun A (a b : F A) => forall (K : forall (A : Set), relation (F A)),
847     (forall (A : Set) (a b : F A), AlgRel R K _ a b) -> K _ a b.
848
849 (a) The algebra and the least fixpoint operators for effects and adverb data types (Alg1, Fix1), and for adverb
850 theories (AlgRel, FixRel).
851 Variant Sum1 (F G : (Set -> Set) -> Set -> Set) K R :=
852   Inl1 (a : F K R) | Inr1 (a : G K R).
853 Variant SumRel {F : Set -> Set}
854   (P Q : (forall (A : Set), relation (F A)) -> forall (A : Set), relation (F A))
855   (K : forall (A : Set), relation (F A)) : forall (A : Set), relation (F A) :=
856   | InlRel {A : Set} {a b : F A} : P K _ a b -> SumRel P Q K _ a b
857   | InrRel {A : Set} {a b : F A} : Q K _ a b -> SumRel P Q K _ a b.
858
859 Notation "F  $\oplus$  G" := (Sum1 F G).
860 Notation "F  $\uplus$  G" := (SumRel F G).

```

(b) The Coq definitions for the \oplus and \uplus operators.

Fig. 10. Key definitions for implementing composable program adverbs in Coq.

The \mathbb{T} data type here is equivalent to the non-composable `ReifiedApp` shown in Fig. 5.

Adverb interpretation can be defined as an algebra of type `Alg1 F E` (Fig. 10a) where `F` is the adverb data type and `E` is the instance we are interpreting to. To apply this “interpretation algebra” to the composed “inductive type”, we fold it over `Fix1` as follows:

```

870 Definition foldFix1 {E A} (alg : Alg1 F E) (f : Fix1 F A) : E A := f _ alg.

```

We define all composable adverb data types using `Set` rather than `Type` because we use the impredicative sets extension in Coq, following MTC. The consequence of this decision is that (1) certain types cannot inhabit in `Set`, and (2) the extension is inconsistent with certain axioms like classical axioms.⁹ We also develop other mechanisms like the injection type classes, the induction principles following MTC. We omit more detail here due to the space constraint. The interested readers can find them in MTC or our supplementary artifact.

Besides MTC, there are other solutions that address the expression problem in theorem provers like Coq. We discuss those alternative solutions in Section 6.

⁹<https://github.com/coq/coq/wiki/Impredicative-Set>

```

883 Variant ReifiedPure (K : Set -> Set) (R : Set) : Set :=
884 | Pure (r : R).
885 Variant ReifiedFunctor (K : Set -> Set) (R : Set) : Set :=
886 | FMap {X : Set} (g : X -> R) (f : K X).
887 Variant ReifiedApp (K : Set -> Set) (R : Set) : Set :=
888 | LiftA2 {X Y : Set} (f : X -> Y -> R) (g : K X) (a : K Y).
889 Variant ReifiedSelective (K : Set -> Set) (R : Set) : Set :=
890 | SelectBy {X Y : Set} (f : X -> ((Y -> R) + R)) (a : K X) (b : K Y).
891 Variant ReifiedMonad (K : Set -> Set) (R : Set) : Set :=
892 | Bind {X : Set} (m : K X) (g : X -> K R).
893
894 (a) The composable adverb data types.
895 Class AppKleenePlus (F : Type -> Type) `{Applicative F} :=
896 { kplus {A} : F A -> F A }.
897 Class FunctorPlus (F : Type -> Type) `{Functor F} :=
898 { plus {A} : F A -> F A -> F A }.
899
900 (* The adverb data type for Repeatedly. *)
901 Variant ReifiedKleenePlus (K : Set -> Set) (R : Set) : Set :=
902 | KPlus : K R -> ReifiedKleenePlus K R.
903 (* The adverb data type for Nondeterministically. *)
904 Variant ReifiedPlus (K : Set -> Set) (R : Set) : Set :=
905 | Plus : K R -> K R -> ReifiedPlus K R.
906

```

(b) The adverb data types of Nondeterministically and Repeatedly.

Fig. 11. The composable adverb data types and add-on adverb data types.

4.3 Add-on Adverbs

Another benefit of making program adverbs composable is that we can now define two add-on adverbs, namely Repeatedly and Nondeterministically, which are not suitable as standalone adverbs. These two adverbs reify two classes of functors, namely AppKleenePlus and FunctorPlus, that we define ourselves. We show these classes of functors and their reifications in Fig. 11b. AppKleenePlus is a subclass of Applicative and represents the “Kleene plus”.¹⁰ It is a “Kleene plus” rather than a “Kleene star” because no empty element is defined. FunctorPlus is similar to the commonly-used Alternative and MonadPlus type classes in Haskell, but contains no empty element and only requires itself to be a subclass of Functor. We define these type classes’ reifications as add-on adverbs so that these adverbs can be composed with classes of functors at different expressive levels: e.g., Repeatedly can be composed with Statically as well as Dynamically.

We show the adverb theories of Repeatedly and Nondeterministically in Fig. 12. Both of these two add-on adverbs are somewhat nondeterministic, so one change we make to their adverb theories is adding refinement relations (\sqsubseteq) in addition to bisimulation relations (\cong).

We show that these two adverbs are sound with respect to the following adverb simulations:

$$\begin{aligned}
 \text{ReifiedKleenePlus} &\models_{\text{PowerSet}} \text{AppKleenePlus} \\
 \text{ReifiedPlus} &\models_{\text{PowerSet}} \text{FunctorPlus}
 \end{aligned}$$

¹⁰https://en.wikipedia.org/wiki/Kleene_star#Kleene_plus

REPEAT	:	$\forall n, \text{repeat } a \ n \subseteq \text{kplus } a$
KPLUS	:	$\frac{a \subseteq \text{kplus } b}{\text{kplus } a \subseteq \text{kplus } b}$
COMMUTATIVITY	:	$\text{plus } a \ b \cong \text{plus } b \ a$
ASSOCIATIVITY	:	$\text{plus } a \ (\text{plus } b \ c) \cong \text{plus } (\text{plus } a \ b) \ c$
PLUS	:	$\frac{a \subseteq c \quad b \subseteq c}{\text{plus } a \ b \subseteq c}$
LEFT PLUS	:	$a \subseteq \text{plus } a \ b$
RIGHT PLUS	:	$b \subseteq \text{plus } a \ b$

Fig. 12. The adverb theories for Repeatedly and Nondeterministically. The function `repeat a n` repeats `a` for `n` times. Functions `kplus` and `plus` are smart constructors of `KPlus` and `Plus`, respectively.

```

(* FunctorPlus transformer. *)
Definition fmapPowerSet {A B : Type} (f : A -> B) (a : PowerSet I A) : PowerSet I B :=
  fun r => exists a', a a' /\ fmap f a' ≡ r.

Definition plusPowerSet {A : Type} (a b : PowerSet I A) : PowerSet I A :=
  fun r => a r \/ b r.

(* AppKleenePlus transformer. *)
Definition liftA2PowerSet {A B C : Type} (f : A -> B -> C)
  (a : PowerSet I A) (b : PowerSet I B) : PowerSet I C :=
  fun r => exists a' b', a a' /\ b b' /\ (liftA2 f a' b' ≡ r).

Fixpoint repeatPowerSet {A : Type} (a : PowerSet I A) (n : nat) : PowerSet I A :=
  match n with
  | 0 => a
  | S n => liftA2PowerSet (fun _ x => x) a (repeatPowerSet a n)
  end.

Definition kplusPowerSet {A : Type} (a : PowerSet I A) : PowerSet I A :=
  fun r => exists n, repeatPowerSet a n r.

```

Fig. 13. The `FunctorPlus` transformer instance and the `AppKleenePlus` transformer instance of the `PowerSet` data type. \equiv is the lawful equivalence relation on original functor/applicative functor `I`.

The definition of `PowerSet` data type is the same as that in Fig. 8, but we are using its `AppKleenePlus` transformer and `FunctorPlus` transformer instances here. The core definitions of these transformers are shown in Fig. 13.

5 EXAMPLES

In this section, we use two different examples to demonstrate the usefulness and different aspects of program adverbs and Tlön embeddings.

5.1 Haxl

In our first example, we show that we can use composable adverbs to capture two different computation patterns in the same library. We also demonstrate interpreting composable adverbs to a shallow embedding in a modular way.

We illustrate these aspects via an example based on the core ideas of Haxl. Haxl is a Haskell library developed and maintained by Meta (formerly known as Facebook) that automatically parallelizes certain operations to achieve better performance [Marlow et al. 2014]. As an example, suppose that we want to fetch data from a database and we have a `Fetch : Type -> Type` data type that encapsulates the fetching effect. The key insight of the Haxl library is to distinguish the operations of `Fetch`'s `Monad` instance and those of its `Applicative` instance. When we use `>>=` to bind two `Fetch`s, those data fetches are sequential; but when we use `liftA2` to bind two of them, those data fetches are batched and will be sent to the database together. To achieve this, it is important that the definition of `liftA2` is not equivalent to the “default” definition derived from `>>=`.

This design of Haxl poses a challenge to mixed embeddings based on freer monads or any other basic adverbs discussed in Section 3, because we need to distinguish when `Applicative` operations are used and when `Monad` operations are used. This is exactly where composable adverbs are useful.

In this example, we assume that we already have a translation from Haxl's `Applicative` and `Monad` operations to those operations in `Coq`.¹¹ In our embedding, we use the following `T` datatype to encode the Tlön embedding of a data fetching program:

Definition `T := Fix1 (ReifiedPure \oplus ReifiedApp \oplus ReifiedMonad \oplus DataEff).`

We use `ReifiedApp` to model batched operations and the theory of `StaticallyInParallel` to model their parallel nature. We use `ReifiedMonad` to model sequential operations.

We cannot know statically how many database accesses would happen in a Haxl program, because a program can choose to do different things depending on the result of some data fetch. Therefore, we need to pick an effect interpretation for `DataEff` to reason about this property. In this example, we are assuming that the database does not change, so we interpret our Tlön embedding to a shallow embedding whose semantic domain is the update monad [Ahman and Uustalu 2013].

The key definitions of the update monad are shown in Fig. 14a. The update monad is essentially a combination of a reader monad and a writer monad. In our example, the “reader state” has type `var -> val` which represents an immutable key-value database we can read from. The “writer state” is a `nat`, which represents the accumulated number of database accesses. The `bind` operation propagates the key-value database and accumulates the cost.

Additionally, we define a `liftA2` operation, which only records the maximum number of database accesses in one of its branches. This is not the same as the `liftA2` operation that can be automatically derived from the monad instance of `Update`. Furthermore, this `liftA2` is commutative. Thanks to that, we can interpret `T` to the `Update` datatype without the help of a `PowerSet` transformer.

Figure 14b shows how we interpret composed adverbs in a modular way. First, we define a type class called `AdverbAlg` for interpretation algebras. We then define an interpretation from each individual composable adverb and effect in `T` to `Update`. Finally, the interpretation from `T` to `Update` can be automatically inferred by `Coq` thanks to the instance `AdverbAlgSum`. If we would like to add another effect or composable adverb to `T`, we only need to add one more instance of `AdverbAlg` and we do not need to modify any existing interpretation algebras.

Interested readers can find the full `Coq` implementation of the `Update` data type, the `AdverbAlg` type class and relevant instances, along with a few simple examples in our supplementary artifact.

¹¹Tools like `hs-to-coq` [Breitner et al. 2021; Spector-Zabusky et al. 2018] can be adapted to implement the translation.

```

1030 Definition Update A := ((var -> val) -> A * nat).
1031
1032 Definition ret {A} (a : A) : Update A := fun map => (a, 0).
1033 Definition bind {A B} (m : Update A) (k : A -> Update B) : Update B :=
1034   fun map => match m map with
1035     | (i, n) => match (k i map) with
1036       | (r, n') => (r, n + n')
1037     end
1038   end.
1039 Definition liftA2 {A B C} (f : A -> B -> C) (a : Update A) (b : Update B) : Update C :=
1040   fun map => match (a map, b map) with
1041     | ((a, n1), (b, n2)) => (f a b, max n1 n2)
1042   end.
1043 Definition get (v : var) : Update val := fun map => (map v, 1).
1044
1045 (a) The Update datatype.
1046 Class AdverbAlg (D : (Set -> Set) -> Set -> Set) (I : Set -> Set) :=
1047   { adverbAlg : Alg1 D I }.
1048
1049 Instance CostApp : AdverbAlg ReifiedApp Update :=
1050   { | adverbAlg := fun d => match d with LiftA2 f a b => liftA2 f a b end | }.
1051 Instance CostMonad : AdverbAlg ReifiedMonad Update :=
1052   { | adverbAlg := fun d => match d with Bind m k => bind m k end | }.
1053 Instance CostPure : AdverbAlg ReifiedPure Update :=
1054   { | adverbAlg := fun d => match d with Pure a => ret a end | }.
1055 Instance CostData : AdverbAlg DataEff Update :=
1056   { | adverbAlg := fun d => match d with GetData v => get v end | }.
1057
1058 Instance AdverbAlgSum D1 D2 I `{AdverbAlg D1 I} `{AdverbAlg D2 I} :
1059   AdverbAlg (D1 ⊕ D2) I name :=
1060   { | adverbAlg := fun a => match a with
1061     | Inl1 a => adverbAlg a
1062     | Inr1 a => adverbAlg a
1063   end | }.
1064

```

(b) Interpretation algebras that interpret composable adverbs and DataEff to Update. Thanks to Instance AdverbAlgSum and Coq's type class inference, we can automatically get the interpretation from T to Update.

Fig. 14. The Update datatype and the interpretation from T to Update.

5.2 A Networked Server

A common technique used in formal verification is dividing the verification into multiple layers and establishing a refinement relation between every two layers [Gu et al. 2018; Koh et al. 2019; Lorch et al. 2020; Zakowski et al. 2021]. This approach offers better abstraction and modularity, as at each layer, we only need to consider certain subsets of properties.

In this example, we show the usefulness of program adverbs and Tlön embeddings in a layered approach. Specifically, we define an *intermediate-level* specification that omits implementation

```

10791 newconn ::<- accept ;;
10802 IF (not (*newconn == 0)) THEN
10813   newconn_rec := connection *newconn
10824   READING ;;
10835   conns :++ newconn_rec
10846 END ;;
10857 FOR y IN conns DO
10868   IF (y->state == WRITING) THEN
10879     r ::<- write y->id *s ;;
10880     y->state := CLOSED
10891   END ;;
10902   IF (y->state == READING) THEN
10913     r ::<- read y->id ;;
10924     IF (*r == 0) THEN
10935       y->state := CLOSED
10946     ELSE
10957       s := *r ;;
10968       y->state := WRITING
10979     END
10980   END
10981 END.

```

```

Some
  (Or (newconn ::<- accept ;;
      IF (not (*newconn == 0)) THEN
        newconn_rec := connection *newconn
        READING ;;
        conns :++ newconn_rec
      END)
    (OneOf (conns) y
      (Or (IF (y->state == WRITING) THEN
          r ::<- write y->id *s ;;
          y->state := CLOSED
        END)
        (IF (y->state == READING) THEN
          r ::<- read y->id ;;
          IF (*r == 0) THEN
            y->state := CLOSED
          ELSE
            s := *r ;;
            y->state := WRITING
          END
        END)))

```

(a) The implementation Impl in NETIMP.

(b) Our intermediate layer specification Spec in NETSPEC.

Fig. 15. The implementation and the intermediate layer specification of our networked server.

details about execution orders, *etc.* Since the specification is only more *nondeterministic* in its control flow, we would like our formal verification to show that an implementation refines the specification *without* interpreting effects to a shallow embedding. This is exactly where program adverbs and Tlön embeddings can help.

We demonstrate this vision above via a simple server adapted from that of Koh et al. [2019]. The server communicates with multiple clients via a network interface. A client initiates a communication with the server by sending a request that is a number. Whenever the server receives a request, it stores the number of that request and sends back a number in its store—a client does not necessarily receive what they send, because the server can interleave multiple sessions.

We show that a specific implementation of such a server refines an intermediate-level specification. We also show the refinements based on Tlön embeddings with the help of adverb theories. Unlike Koh et al. [2019], we do not show that the implementation further refines a higher-level specification based on observations over a network, as that is beyond the scope of this work.

The implementation. The server is implemented using a single-process *event loop* [Pai et al. 1999]. Instead of processing a request and sending back a response immediately, the server divides a session with a client into multiple steps. In each iteration of the event loop, the server advances the session of each request by one step, thus interleaving multiple sessions.

We show the main loop body of our adapted version of the networked server in Fig. 15a. For simplicity, we use a custom language called NETIMP. NETIMP supports datatypes like booleans, natural numbers, and a special record type called connection. It has network operations like accept, read, and write. All these operations return natural numbers, with 0 indicating failures. The language does not have a while loop but it has a FOR loop that iterates over a list. The loop

Definition L1 := A ;; FOR y IN conns DO B ;; C ;; END.	Definition L2 := A ;; OneOf (conns) y (B ;; C).	Definition L3 := A ;; OneOf (conns) y (Or B C).
--	---	---

Fig. 16. Program L1 written in NETIMP, and programs L2 and L3 written in NETSPEC.

variable is implemented as a pointer that points to elements in the list iteratively. We also use C-like notations (*i.e.*, * and ->) for operations on pointers.

The implementation `Impl` maintains a list of connections called `conns`. Each connection in `conns` is in one of the three possible states: `READING`, `WRITING`, or `CLOSED`. At the start of each loop, the server checks if there is a new connection waiting to be established by calling the non-blocking operation `accept`. If there is, the server creates a new connection with the `READING` state and adds it to `conns`. The server then goes over each connection in `conns`: if a connection is in the `READING` state, the server tries to read from the connection and updates an internal state `s` with the recently read value; if a connection is in the `WRITING` state, the server sends the current value of its internal state `s` to the connection; once a connection enters the `CLOSED` state, it remains that state forever and the server will not do anything with it—we design the server in this way for simplicity; a more realistic server should remove closed connections from `conns`.

The specification. We show our specification `Spec` in Fig. 15b. `Spec` is written in a language called `NETSPEC`. `NETSPEC` adds a few additional commands to `NETIMP`: `Some` is a unary operation that models the “Kleene plus”; `Or` is a binary operation that models a nondeterministic choice wrapped inside a “Kleene plus”; `OneOf` is like `Or`, but it nondeterministically chooses from a list—line 8 means that we nondeterministically assign the variable `y` with one element from the list in `conns`.

`Spec` is more nondeterministic compared with `Impl`. At each iteration of the event loop, `Impl` always first tries to accept a connection. It then goes over the list of `conns` in a fixed order. `Spec` does not enforce order: an `accept` could happen immediately after another `accept`; we can access elements in `conns` in any order and some connection might get visited more often than others.

Tlön embeddings and the refinement proof. To show that `Impl` refines `Spec`, we embed both `NETIMP` and `NETSPEC` in Coq using program adverbs. We use the following datatype:

Definition $T := \text{Fix1 } (\text{ReifiedKleenePlus} \oplus \text{ReifiedPlus} \oplus \text{ReifiedPure} \oplus \text{ReifiedMonad} \oplus \text{NetworkEff} \oplus \text{MemoryEff} \oplus \text{FailEff}).$

We have already seen the first four adverbs. Effect `NetworkEff` models the effects incurred by network operations `accept`, `read`, and `write`. Effect `MemoryEff` models the effects incurred by assigning values to variables and retrieving values from them. Finally, effect `FailEff` models when the program crashes.

We use $\llbracket \cdot \rrbracket_T^I$ to denote `NETIMP`’s Tlön embedding in T and $\llbracket \cdot \rrbracket_T^S$ to denote `NETSPEC`’s Tlön embedding in T . For the sake of space, we omit the embeddings here.

We would like to show that $\llbracket \text{Impl} \rrbracket_T^I \subseteq \llbracket \text{Spec} \rrbracket_T^S$. Recall that \subseteq is the refinement relation on program adverbs (Section 4.3). The theorem states that the Tlön embedding of our implementation `Impl` in T refines the Tlön embedding of our specification `Spec` in T .

To show that, we first observe that `Impl` and `Spec` share some common program fragments, *e.g.*, lines 1–6 of `Impl` are the same as lines 2–7 of `Spec`. Indeed, there are three such common

fragments and we name them A (lines 1–6 of `Impl`), B (lines 8–11 of `Impl`), and C (lines 12–20 of `Impl`), respectively. We then define three programs $L1$, $L2$, and $L3$ shown in Fig. 16. These programs represent some intermediate layers between `Impl` and `Spec`. We prove the following theorem:

THEOREM 5.1. $\llbracket \text{Impl} \rrbracket_T^I \subseteq \llbracket L1 \rrbracket_T^I \subseteq \llbracket L2 \rrbracket_T^S \subseteq \llbracket L3 \rrbracket_T^S \subseteq \llbracket \text{Spec} \rrbracket_T^S$.

PROOF. We show $\llbracket \text{Impl} \rrbracket_T^I \subseteq \llbracket L1 \rrbracket_T^I$ by associativity of `Dynamically`. Both $\llbracket L1 \rrbracket_T^I \subseteq \llbracket L2 \rrbracket_T^S$ and $\llbracket L2 \rrbracket_T^S \subseteq \llbracket L3 \rrbracket_T^S$ can be proven by an induction over conns and with the help of theories of `Dynamically`, `Repeatedly` and `Nondeterministically`. Finally, we prove $\llbracket L3 \rrbracket_T^S \subseteq \llbracket \text{Spec} \rrbracket_T^S$ by the theories of `Dynamically`, `Repeatedly`, and `Nondeterministically`. \square

Interested readers can find the full Coq implementation of `NETIMP`, `NETSPEC`, the Tlön embeddings of these two languages, the implementation `Impl`, the specification `Spec`, as well as the full proof of Theorem 5.1 in our supplementary artifact.

6 DISCUSSION

The expression problem. The composable program adverbs require extensible inductive types. We implement this feature in Coq by using the Church encodings of datatypes, following the precedent work of MTC [Delaware et al. 2013]. There are several consequences of using Church encodings instead of Coq’s original inductive datatypes.

First, we cannot make use of Coq’s language mechanisms, libraries, and plugins that make use of Coq’s inductive types (e.g., Coq’s builtin induction principle generator, the `Equations` plugin [Sozeau and Mangin 2019], the `QuickChick` plugin [Lampropoulos et al. 2018; Paraskevopoulou et al. 2022], etc.). Furthermore, the extra implementation overheads incurred by Church encodings (e.g., proving an algebra is a functor, proving the induction principle using dependent types, etc.) can be huge. However, this situation can be helped by developing tools or plugins for supporting Church encodings.

The other consequence is that, following the practice of MTC, we use Coq’s impredicative set extension. This causes (1) certain types cannot inhabit in `Set`, and (2) our Coq development to be inconsistent with certain axioms like classical axioms, as we have discussed in Section 4.2.

There are alternative methods for addressing the expression problem. One option is the meta-programming approach proposed by Forster and Stark [2020]. In this approach, we can define each composable adverb separately in a meta language and use a language plugin to generate a combined definition in Coq. This approach does not fully address the expression problem as extending the combined definition requires recompilation—but the amount of code that needs to be recompiled is much smaller and the generated code uses Coq’s builtin inductive types. Another option that has recently been explored by Kravchuk-Kirilyuk et al. [2021] is adding *family polymorphism* [Ernst 2001] to theorem provers. These works are promising. Unfortunately, they either lack mature tool support or is still in development at the moment. We would like to explore these approaches in the future and composable program adverbs might provide a good application to these approaches.

Reified vs. free structures. Even though the reified structures used in adverb data types are free structures, they are different from those free structures present in Capriotti and Kaposi [2014]; Kiselyov and Ishii [2015]; Mokhov [2019]; Mokhov et al. [2019]. The biggest difference between reified structures and these free structures are the parameters they recurse on: all the reified structures recurse on both their computational parameters, while each free structure only recurses on one of them.¹² For example, comparing `FreerMonad` in Fig. 2 and `ReifiedMonad` in Fig. 5: `FreerMonad` only recurses on the parameter `k` of `Bind`, while `ReifiedMonad` recurses on both

¹²With the exception of reified/free functors, since each of them has only one computational parameters to be recursed on.

parameters m and k . This means that a free structure does not just reify a class of functors, it also converts the reification to a left- or right-associative normal form.

One advantage of the normal forms in free structure definitions is that the type class laws can be automatically derived from definitional equality (with the help of the axiom of functional extensionality). However, this conversion would eliminate some differences in the syntax. Taking `ReifiedApp` as an example, normalizing it would result in a “list” rather than a “binary tree”, making analyzing the depth of the tree impossible. Preserving the original tree structure of `StaticallyInParallel` also plays a crucial role in our examples shown in [Section 2.4](#) and [5.1](#).

7 RELATED WORK

Semantic embeddings. There are various works that study different semantic embeddings. [Boulton et al. \[1992\]](#) are the pioneers who coined terms such as semantic embeddings, shallow embeddings, and deep embeddings. It is known that there are many styles of embeddings between shallow and deep embeddings, but there is not an agreed term on describing them. In this paper, we use the term *mixed embeddings*, which is borrowed from [Chlipala \[2021\]](#), where it is used to describe an embedding based on freer monads. Another term *deeper shallow embeddings* is proposed by [Prinz et al. \[2022\]](#), which shows a way of deepening any shallow embedding.

Freer Monads and Variants. Freer monads [\[Kiselyov and Ishii 2015\]](#) and their variants are studied by many researchers in formal verification to reason about programs with effects. For example, [Letan et al. \[2021\]](#) use freer monads to develop a modular verification framework based on effects and effect handlers called `FreeSpec`. [Christiansen et al. \[2019\]](#) develop a framework based on free monads and containers for reasoning about Haskell programs with effects. [Swierstra and Baanen \[2019\]](#) interpret freer monads into a predicate transformer semantics that is similar to Dijkstra monads; [Nigron and Dagand \[2021\]](#) interprets freer monads using separation logic.

On the *coinductive* side, [Xia et al. \[2020\]](#) develop a coinductive variant of freer monads called *interaction trees* that can be used to reason about general recursions and nonterminating programs. [Koh et al. \[2019\]](#) encode interaction trees in VST [\[Appel et al. 2014\]](#) to reason about networked servers. [Mansky et al. \[2020\]](#) use interaction trees as a lingua franca to interface and compose higher-order separation logic in VST and a first-order verified operating system called `CertiKOS` [\[Gu et al. 2019\]](#). [Zakowski et al. \[2020\]](#) propose a technique called generalized parameterized coinduction for developing equational theory for reasoning about interaction trees. [Zakowski et al. \[2021\]](#) use interaction trees to define a modular, compositional, and executable semantics for LLVM. [Yoon et al. \[2022\]](#) further extend the modularity of interaction trees by extending them with layered monadic interpreters. [Silver and Zdancewic \[2021\]](#) connect interaction trees with Dijkstra monads [\[Maillard et al. 2019\]](#) for writing termination sensitive specifications based on uninterpreted effects. [Lesani et al. \[2022\]](#) use interaction trees to verify transactional objects.

Among many variants of freer monads, one particular structure closely resembles program adverbs. That is the action trees defined in [Swamy et al. \[2020\]](#). The action trees have four constructors, `Act`, `Ret`, `Par`, and `Bind`, whose types correspond to effects, `ReifiedPure`, `ReifiedApp`, and `ReifiedMonad` in composable program adverbs, respectively, another evidence that program adverbs are general models. In contrast to our work, compositionality and extensibility of “adverbs” are not the main issue action trees try to address, so action trees are not built in a composable way. On the other hand, action trees are embedded with separation logic assertions, which are not the focus of `Tlön` embeddings or program adverbs.

Other Free Structures. Other free structures are also explored by various works. [Capriotti and Kaposi \[2014\]](#) propose two variants of freer applicative functors, which correspond to the left- and right-associative variants, respectively. [Xia \[2019\]](#) explores defining freer applicative functors

in Coq, and points out that the right associative variant is harder to define in Coq. Milewski [2018] discusses how to derive freer monoidal functors.¹³ Mokhov [2019] defines the freer selective applicative functors.

Programming Abstractions. We are not the first to observe that monads are too dynamic for certain applications. For example, Swierstra and Duponcheel [1996] identify that a parser that has some static features cannot be defined as a monad. Inspired by their observation, Hughes [2000] proposes a new abstract interface called arrows. The relationship among arrows, applicative functors, monads are studied by Lindley et al. [2011]. Willis et al. [2020] observe that monads generate dynamic structures that are hard to optimize. They further show that, by using applicative and selective functors instead, it is possible to implement staged parser combinators that generate efficient parsers. Mokhov et al. [2020] observe that the datatype of tasks in a build system (called Task in their paper) can be parameterized by a class constraint to describe various kinds of build tasks. For example, a Task Applicative describes tasks whose dependencies are determined *statically* without running the task; and a Task Monad describes tasks with *dynamic* dependencies.

8 CONCLUSION

In this paper, we compare different styles of semantic embeddings and how they impact formal reasoning about programs with effects. We find that, if used properly, mixed embeddings can combine benefits of both shallow and deep embeddings, and be effective in (1) preserving syntactic structures of original programs, (2) showing general properties that can be proved without assumptions over external environment, and (3) reasoning about properties in specialized semantic domains.

We propose *program adverbs* and *Tlön embeddings*, a class of structures and a style of mixed embeddings based on these structures, that enable us to reap these benefits. Like free monads, program adverbs embed pure computations shallowly and effects deeply (and abstractly, but can later be interpreted). However, various program adverbs correspond to alternative computation patterns, and can be composed to model programs with multiple characteristics.

Based on program adverbs, Tlön embeddings cover a wide range of programs and allow us to reason about syntactic properties, semantic properties, and general semantic properties with no assumption over external environment within the same embedding.

ACKNOWLEDGMENT

We thank the anonymous ICFP’22 reviewers, as well as the POPL’22 and ESOP’22 reviewers, whose valuable feedback helped improve this paper. We thank various researchers for their feedback during their discussions with the authors, including Stephanie Balzer, Conal Elliott, Yannick Forster, Paul He, Apoorv Ingle, Ende Jin, Konstantinos Kallas, Anastasiya Kravchuk-Kirilyuk, Andrey Mokhov, Benjamin Pierce, Nick Rioux, Antal Spector-Zabusky, Caleb Stanford, Kathrin Stark, Nikhil Swamy, Val Tannen, Steve Zdancewic, Weixin Zhang, and Yizhou Zhang, *etc.* We thank the anonymous ICFP’22 artifact reviewers for their comments and suggestions that helped improve the artifact. We thank Paolo Giarrusso (on CoqClub Zulipchat), Li-yao Xia, and Irene Yoon for helping with some Coq issues the authors encountered developing the artifact.

This material is based upon work supported by the National Science Foundation under Grant No. 1521539, Grant No. 1703835, and Grant No. 2006535. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the National Science Foundation.

¹³Monoidal functors are equivalent to applicative functors, so they also correspond to the *Statically* adverb.

REFERENCES

- Danel Ahman and Tarmo Uustalu. 2013. Update Monads: Cointerpreting Directed Containers. In *19th International Conference on Types for Proofs and Programs, TYPES 2013, April 22-26, 2013, Toulouse, France (LIPIcs, Vol. 26)*, Ralph Matthes and Aleksy Schubert (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 1–23. <https://doi.org/10.4230/LIPIcs.TYPES.2013.1>
- Andrew W. Appel, Robert Dockins, Aquinas Hobor, Lennart Beringer, Josiah Dodds, Gordon Stewart, Sandrine Blazy, and Xavier Leroy. 2014. *Program Logics for Certified Compilers*. Cambridge University Press. <http://www.cambridge.org/de/academic/subjects/computer-science/programming-languages-and-applied-logic/program-logics-certified-compilers?format=HB>
- Brian E. Aydemir, Aaron Bohannon, Matthew Fairbairn, J. Nathan Foster, Benjamin C. Pierce, Peter Sewell, Dimitrios Vytiniotis, Geoffrey Washburn, Stephanie Weirich, and Steve Zdancewic. 2005. Mechanized Metatheory for the Masses: The POPLMARK Challenge. In *Theorem Proving in Higher Order Logics, 18th International Conference, TPHOLs 2005, Oxford, UK, August 22-25, 2005, Proceedings*. 50–65. https://doi.org/10.1007/11541868_4
- Jorge Luis Borges. 1940. Tlön, Uqbar, Orbis Tertius. In *Labyrinths*, Donald A. Yates and James E. Irby (Eds.). New Directions.
- Richard J. Boulton, Andrew D. Gordon, Michael J. C. Gordon, John Harrison, John Herbert, and John Van Tassel. 1992. Experience with Embedding Hardware Description Languages in HOL. In *Theorem Provers in Circuit Design, Proceedings of the IFIP TC10/WG 10.2 International Conference on Theorem Provers in Circuit Design: Theory, Practice and Experience, Nijmegen, The Netherlands, 22-24 June 1992, Proceedings (IFIP Transactions, Vol. A-10)*, Victoria Stavridou, Thomas F. Melham, and Raymond T. Boute (Eds.). North-Holland, 129–156.
- Joachim Breitner, Antal Spector-Zabusky, Yao Li, Christine Rizkallah, John Wiegley, Joshua Cohen, and Stephanie Weirich. 2021. Ready, Set, Verify! Applying hs-to-coq to real-world Haskell code. *Journal of Functional Programming* 31 (2021), e5. <https://doi.org/10.1017/S0956796820000283>
- Paolo Capriotti and Ambrus Kaposi. 2014. Free Applicative Functors. In *Proceedings 5th Workshop on Mathematically Structured Functional Programming, MSFP@ETAPS 2014, Grenoble, France, 12 April 2014 (EPTCS, Vol. 153)*, Paul Levy and Neel Krishnaswami (Eds.). 2–30. <https://doi.org/10.4204/EPTCS.153.2>
- Adam Chlipala. 2021. Skipping the binder bureaucracy with mixed embeddings in a semantics course (functional pearl). *Proc. ACM Program. Lang.* 5, ICFP (2021), 1–28. <https://doi.org/10.1145/3473599>
- Jan Christiansen, Sandra Dylus, and Niels Bunkenburg. 2019. Verifying effectful Haskell programs in Coq. In *Proceedings of the 12th ACM SIGPLAN International Symposium on Haskell, Haskell@ICFP 2019, Berlin, Germany, August 18-23, 2019*, Richard A. Eisenberg (Ed.). ACM, 125–138. <https://doi.org/10.1145/3331545.3342592>
- Bruno C. d. S. Oliveira. 2009. Modular Visitor Components. In *ECOOP 2009 - Object-Oriented Programming, 23rd European Conference, Genoa, Italy, July 6-10, 2009, Proceedings (Lecture Notes in Computer Science, Vol. 5653)*, Sophia Drossopoulou (Ed.). Springer, 269–293. https://doi.org/10.1007/978-3-642-03013-0_13
- Benjamin Delaware, Bruno C. d. S. Oliveira, and Tom Schrijvers. 2013. Meta-theory à la carte. In *The 40th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '13, Rome, Italy - January 23 - 25, 2013*, Roberto Giacobazzi and Radhia Cousot (Eds.). ACM, 207–218. <https://doi.org/10.1145/2429069.2429094>
- Sandra Dylus, Jan Christiansen, and Finn Teegen. 2019. One Monad to Prove Them All. *Art Sci. Eng. Program.* 3, 3 (2019), 8. <https://doi.org/10.22152/programming-journal.org/2019/3/8>
- Erik Ernst. 2001. Family Polymorphism. In *ECOOP 2001 - Object-Oriented Programming, 15th European Conference, Budapest, Hungary, June 18-22, 2001, Proceedings (Lecture Notes in Computer Science, Vol. 2072)*, Jørgen Lindskov Knudsen (Ed.). Springer, 303–326. https://doi.org/10.1007/3-540-45337-7_17
- Yannick Forster and Kathrin Stark. 2020. Coq à la carte: a practical approach to modular syntax with binders. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2020, New Orleans, LA, USA, January 20-21, 2020*, Jasmin Blanchette and Catalin Hritcu (Eds.). ACM, 186–200. <https://doi.org/10.1145/3372885.3373817>
- Simon Foster, Chung-Kil Hur, and Jim Woodcock. 2021. Formally Verified Simulations of State-Rich Processes using Interaction Trees in Isabelle/HOL. *CoRR* abs/2105.05133 (2021). arXiv:2105.05133 <https://arxiv.org/abs/2105.05133>
- Ronghui Gu, Zhong Shao, Hao Chen, Jieung Kim, Jérémie Koenig, Xiongnan (Newman) Wu, Vilhelm Sjöberg, and David Costanzo. 2019. Building certified concurrent OS kernels. *Commun. ACM* 62, 10 (2019), 89–99. <https://doi.org/10.1145/3356903>
- Ronghui Gu, Zhong Shao, Jieung Kim, Xiongnan (Newman) Wu, Jérémie Koenig, Vilhelm Sjöberg, Hao Chen, David Costanzo, and Tahina Ramanandro. 2018. Certified concurrent abstraction layers. In *Proceedings of the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2018, Philadelphia, PA, USA, June 18-22, 2018*, Jeffrey S. Foster and Dan Grossman (Eds.). ACM, 646–661. <https://doi.org/10.1145/3192366.3192381>
- John Hughes. 2000. Generalising monads to arrows. *Sci. Comput. Program.* 37, 1-3 (2000), 67–111. [https://doi.org/10.1016/S0167-6423\(99\)00023-4](https://doi.org/10.1016/S0167-6423(99)00023-4)
- Oleg Kiselyov and Hiromi Ishii. 2015. Freer monads, more extensible effects. In *Proceedings of the 8th ACM SIGPLAN Symposium on Haskell, Haskell 2015, Vancouver, BC, Canada, September 3-4, 2015*. 94–105. <https://doi.org/10.1145/2804302.2804319>

- Nicolas Koh, Yao Li, Yishuai Li, Li-yao Xia, Lennart Beringer, Wolf Honoré, William Mansky, Benjamin C. Pierce, and Steve Zdancewic. 2019. From C to interaction trees: specifying, verifying, and testing a networked server. In *Proceedings of the 8th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2019, Cascais, Portugal, January 14-15, 2019*, Assia Mahboubi and Magnus O. Myreen (Eds.). ACM, 234–248. <https://doi.org/10.1145/3293880.3294106>
- Anastasiya Kravchuk-Kirilyuk, Yizhou Zhang, and Nada Amin. 2021. Family Polymorphism for Proof Extensibility. In *Proceedings of the 27th International Conference on Types for Proofs and Programs, TYPES 2021, June 14–18, 2021*. <https://types21.liacs.nl/download/family-polymorphism-for-proof-extensibility/>
- Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C. Pierce. 2018. Generating good generators for inductive relations. *Proc. ACM Program. Lang.* 2, POPL (2018), 45:1–45:30. <https://doi.org/10.1145/3158133>
- Mohsen Lesani, Li-yao Xia, Anders Kaseorg, Christian J. Bell, Adam Chlipala, Benjamin C. Pierce, and Steve Zdancewic. 2022. C4: verified transactional objects. *Proc. ACM Program. Lang.* 6, OOPSLA (2022), 1–31. <https://doi.org/10.1145/3527324>
- Thomas Letan, Yann Régis-Gianas, Pierre Chifflier, and Guillaume Hiet. 2021. Modular verification of programs with effects and effects handlers. *Formal Aspects Comput.* 33, 1 (2021), 127–150. <https://doi.org/10.1007/s00165-020-00523-2>
- Yao Li and Stephanie Weirich. 2022. *Program Adverbs and Tlön Embeddings (artifact)*. <https://doi.org/10.5281/zenodo.6633086>
- Sam Lindley, Philip Wadler, and Jeremy Yallop. 2011. Idioms are Oblivious, Arrows are Meticulous, Monads are Promiscuous. *Electron. Notes Theor. Comput. Sci.* 229, 5 (2011), 97–117. <https://doi.org/10.1016/j.entcs.2011.02.018>
- Jacob R. Lorch, Yixuan Chen, Manos Kapritsos, Bryan Parno, Shaz Qadeer, Upamanyu Sharma, James R. Wilcox, and Xueyuan Zhao. 2020. Armada: low-effort verification of high-performance concurrent programs. In *Proceedings of the 41st ACM SIGPLAN International Conference on Programming Language Design and Implementation, PLDI 2020, London, UK, June 15-20, 2020*, Alastair F. Donaldson and Emina Torlak (Eds.). ACM, 197–210. <https://doi.org/10.1145/3385412.3385971>
- Kenji Maillard, Danel Ahman, Robert Atkey, Guido Martínez, Catalin Hritcu, Exequiel Rivas, and Éric Tanter. 2019. Dijkstra monads for all. *Proc. ACM Program. Lang.* 3, ICFP (2019), 104:1–104:29. <https://doi.org/10.1145/3341708>
- William Mansky, Wolf Honoré, and Andrew W. Appel. 2020. Connecting Higher-Order Separation Logic to a First-Order Outside World. In *Programming Languages and Systems - 29th European Symposium on Programming, ESOP 2020, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2020, Dublin, Ireland, April 25-30, 2020, Proceedings (Lecture Notes in Computer Science, Vol. 12075)*, Peter Müller (Ed.). Springer, 428–455. https://doi.org/10.1007/978-3-030-44914-8_16
- Simon Marlow, Louis Brandy, Jonathan Coens, and Jon Purdy. 2014. There is no fork: an abstraction for efficient, concurrent, and concise data access. In *Proceedings of the 19th ACM SIGPLAN international conference on Functional programming, Gothenburg, Sweden, September 1-3, 2014*, Johan Jeuring and Manuel M. T. Chakravarty (Eds.). ACM, 325–337. <https://doi.org/10.1145/2628136.2628144>
- Coq development team. 2022. *The Coq proof assistant*. <http://coq.inria.fr> Version 8.15.1.
- Conor McBride. 2015. Turing-Completeness Totally Free. In *Mathematics of Program Construction - 12th International Conference, MPC 2015, Königswinter, Germany, June 29 - July 1, 2015. Proceedings*. 257–275. https://doi.org/10.1007/978-3-319-19797-5_13
- Conor McBride and Ross Paterson. 2008. Applicative programming with effects. *J. Funct. Program.* 18, 1 (2008), 1–13. <https://doi.org/10.1017/S0956796807006326>
- Bartosz Milewski. 2018. Free Monoidal Functors, Categorically! <https://bartoszmilewski.com/2018/05/16/free-monoidal-functors-categorically/> Blog post.
- Eugenio Moggi. 1991. Notions of Computation and Monads. *Inf. Comput.* 93, 1 (1991), 55–92. [https://doi.org/10.1016/0890-5401\(91\)90052-4](https://doi.org/10.1016/0890-5401(91)90052-4)
- Andrey Mokhov. 2019. Implementation of selective applicative functors in Haskell. <https://hackage.haskell.org/package/selective>
- Andrey Mokhov, Georgy Lukyanov, Simon Marlow, and Jérémie Dimino. 2019. Selective applicative functors. *Proc. ACM Program. Lang.* 3, ICFP (2019), 90:1–90:29. <https://doi.org/10.1145/3341694>
- Andrey Mokhov, Neil Mitchell, and Simon Peyton Jones. 2020. Build systems à la carte: Theory and practice. *J. Funct. Program.* 30 (2020), e11. <https://doi.org/10.1017/S0956796820000088>
- Pierre Nigron and Pierre-Évariste Dagand. 2021. Reaching for the Star: Tale of a Monad in Coq. In *12th International Conference on Interactive Theorem Proving, ITP 2021, June 29 to July 1, 2021, Rome, Italy (Virtual Conference) (LIPIcs, Vol. 193)*, Liron Cohen and Cezary Kaliszyk (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 29:1–29:19. <https://doi.org/10.4230/LIPIcs.ITP.2021.29>
- Vivek S. Pai, Peter Druschel, and Willy Zwaenepoel. 1999. Flash: An efficient and portable Web server. In *Proceedings of the 1999 USENIX Annual Technical Conference, June 6-11, 1999, Monterey, California, USA*. USENIX, 199–212. http://www.usenix.org/events/usenix99/full_papers/pai/pai.pdf
- Zoe Paraskevopoulou, Aaron Eline, and Leonidas Lampropoulos. 2022. Computing correctly with inductive relations. In *PLDI '22: 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation, San Diego, CA, USA, June 13 - 17, 2022*, Ranjit Jhala and Isil Dillig (Eds.). ACM, 966–980. <https://doi.org/10.1145/3519939.3523707>

- Jacob Prinz, G. A. Kavvos, and Leonidas Lampropoulos. 2022. Deeper Shallow Embeddings. In *13th International Conference on Interactive Theorem Proving, ITP 2022, August 7 to August 10, 2022, Haifa, Israel (LIPIcs)*. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 19:1–19:18. <https://doi.org/10.4230/LIPIcs.ITP.2022.19>
- Lucas Silver and Steve Zdancewic. 2021. Dijkstra monads forever: termination-sensitive specifications for interaction trees. *Proc. ACM Program. Lang.* 5, POPL (2021), 1–28. <https://doi.org/10.1145/3434307>
- Matthieu Sozeau and Cyprien Mangin. 2019. Equations reloaded: high-level dependently-typed functional programming and proving in Coq. *Proc. ACM Program. Lang.* 3, ICFP (2019), 86:1–86:29. <https://doi.org/10.1145/3341690>
- Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. 2018. Total Haskell is reasonable Coq. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los Angeles, CA, USA, January 8–9, 2018*, June Andronick and Amy P. Felty (Eds.). ACM, 14–27. <https://doi.org/10.1145/3167092>
- Josef Svenningsson and Emil Axelsson. 2012. Combining Deep and Shallow Embedding for EDSL. In *Trends in Functional Programming - 13th International Symposium, TFP 2012, St. Andrews, UK, June 12–14, 2012, Revised Selected Papers (Lecture Notes in Computer Science, Vol. 7829)*, Hans-Wolfgang Loidl and Ricardo Peña (Eds.). Springer, 21–36. https://doi.org/10.1007/978-3-642-40447-4_2
- Nikhil Swamy, Aseem Rastogi, Aymeric Fromherz, Denis Merigoux, Danel Ahman, and Guido Martínez. 2020. SteelCore: an extensible concurrent separation logic for effectful dependently typed programs. *Proc. ACM Program. Lang.* 4, ICFP (2020), 121:1–121:30. <https://doi.org/10.1145/3409003>
- Nikhil Swamy, Joel Weinberger, Cole Schlesinger, Juan Chen, and Benjamin Livshits. 2013. Verifying higher-order programs with the dijkstra monad. In *ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '13, Seattle, WA, USA, June 16–19, 2013*, Hans-Juergen Boehm and Cormac Flanagan (Eds.). ACM, 387–398. <https://doi.org/10.1145/2491956.2491978>
- S. Doaitse Swierstra and Luc Duponcheel. 1996. Deterministic, Error-Correcting Combinator Parsers. In *Advanced Functional Programming, Second International School, Olympia, WA, USA, August 26–30, 1996, Tutorial Text (Lecture Notes in Computer Science, Vol. 1129)*, John Launchbury, Erik Meijer, and Tim Sheard (Eds.). Springer, 184–207. https://doi.org/10.1007/3-540-61628-4_7
- Wouter Swierstra and Tim Baanen. 2019. A predicate transformer semantics for effects (functional pearl). *Proc. ACM Program. Lang.* 3, ICFP (2019), 103:1–103:26. <https://doi.org/10.1145/3341707>
- Philip Wadler. 1990. Recursive types for free! <http://homepages.inf.ed.ac.uk/wadler/papers/free-rectypes/free-rectypes.txt> Draft.
- Philip Wadler. 1992. Comprehending Monads. *Math. Struct. Comput. Sci.* 2, 4 (1992), 461–493. <https://doi.org/10.1017/S0960129500001560>
- Philip Wadler. 1998. The Expression Problem. <http://homepages.inf.ed.ac.uk/wadler/papers/expression/expression.txt> Email correspondence.
- Jamie Willis, Nicolas Wu, and Matthew Pickering. 2020. Staged selective parser combinators. *Proc. ACM Program. Lang.* 4, ICFP (2020), 120:1–120:30. <https://doi.org/10.1145/3409002>
- Li-yao Xia. 2019. Free applicative functors in Coq. <https://blog.poisson.chat/posts/2019-07-14-free-applicative-functors.html> Blog post.
- Li-yao Xia, Yannick Zakowski, Paul He, Chung-Kil Hur, Gregory Malecha, Benjamin C. Pierce, and Steve Zdancewic. 2020. Interaction trees: representing recursive and impure programs in Coq. *Proc. ACM Program. Lang.* 4, POPL (2020), 51:1–51:32. <https://doi.org/10.1145/3371119>
- Irene Yoon, Yannick Zakowski, and Steve Zdancewic. 2022. Formal reasoning about layered monadic interpreters. *Proc. ACM Program. Lang.* 6, ICFP (2022).
- Yannick Zakowski, Calvin Beck, Irene Yoon, Ilia Zaichuk, Vadim Zaliva, and Steve Zdancewic. 2021. Modular, compositional, and executable formal semantics for LLVM IR. *Proc. ACM Program. Lang.* 5, ICFP (2021), 1–30. <https://doi.org/10.1145/3473572>
- Yannick Zakowski, Paul He, Chung-Kil Hur, and Steve Zdancewic. 2020. An equational theory for weak bisimulation via generalized parameterized coinduction. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2020, New Orleans, LA, USA, January 20–21, 2020*, Jasmin Blanchette and Catalin Hritcu (Eds.). ACM, 71–84. <https://doi.org/10.1145/3372885.3373813>
- Hengchu Zhang, Wolf Honoré, Nicolas Koh, Yao Li, Yishuai Li, Li-yao Xia, Lennart Beringer, William Mansky, Benjamin C. Pierce, and Steve Zdancewic. 2021. Verifying an HTTP Key-Value Server with Interaction Trees and VST. In *12th International Conference on Interactive Theorem Proving, ITP 2021, June 29 to July 1, 2021, Rome, Italy (Virtual Conference) (LIPIcs, Vol. 193)*, Liron Cohen and Cezary Kaliszyk (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 32:1–32:19. <https://doi.org/10.4230/LIPIcs.ITP.2021.32>